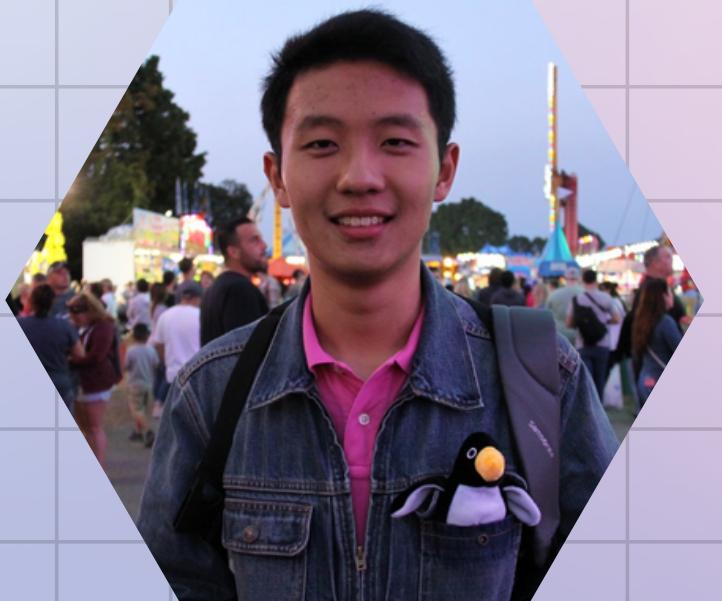




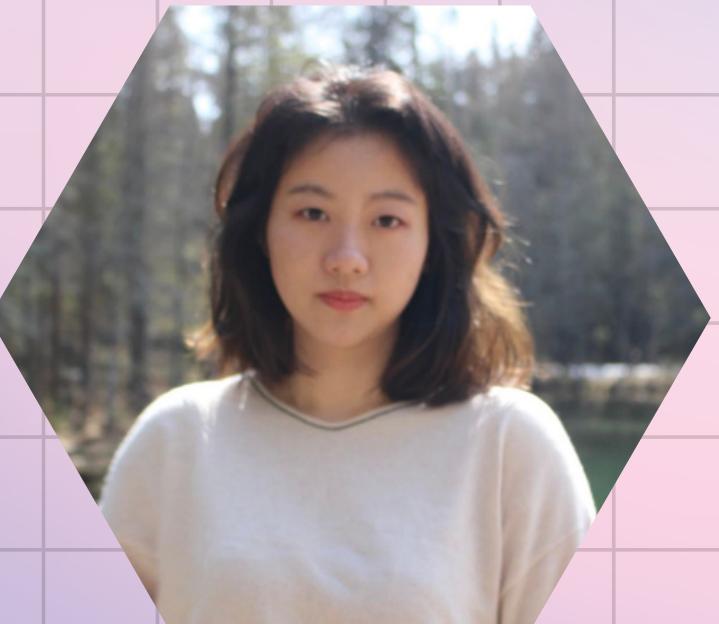
Team 15: Jeffery (Yan) Chen, Olivia (Caixin) Yang, Wilson McDermott, Zhizhuo Li

Our Team

o o o o



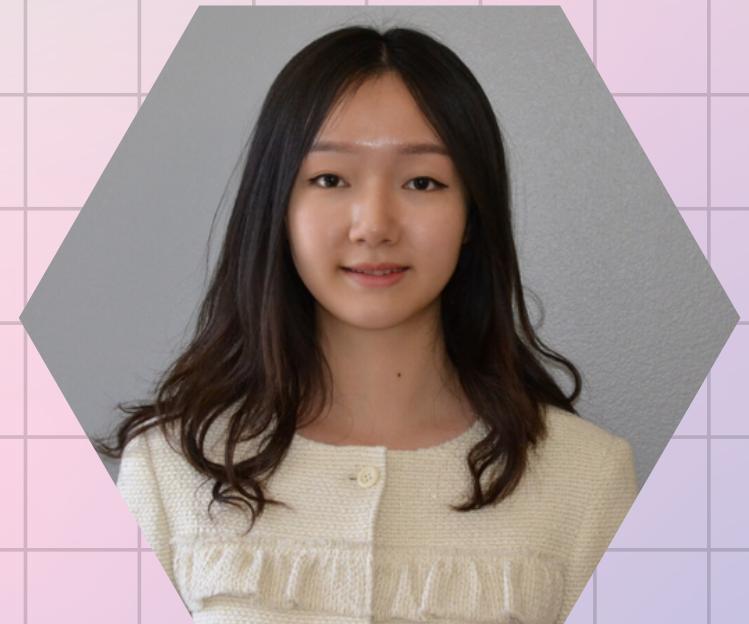
o o o o



o o o o



o o o o



o o o o

Jeffrey Chen

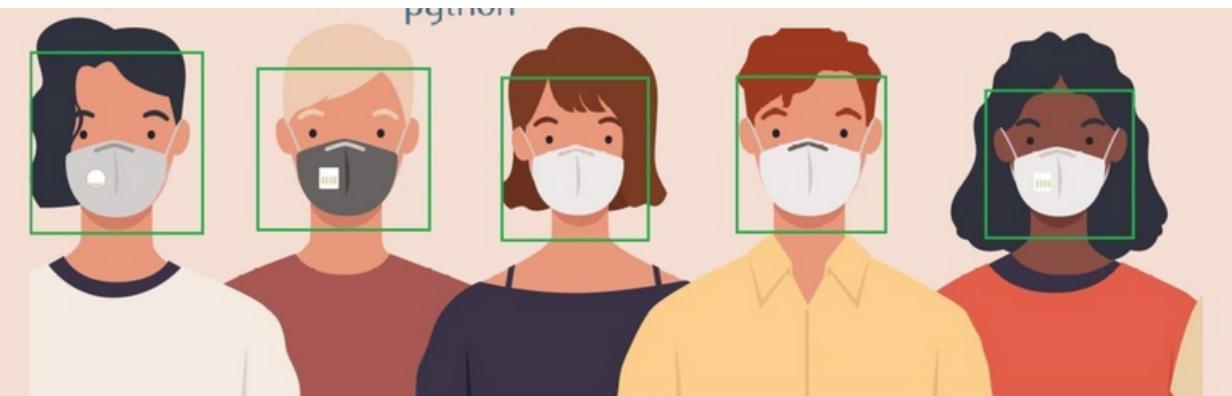
Zhizhuo Li

Wilson McDermott

Caixin (Olivia) Yang

Today's Agenda

o o o o



- 1 Problem Statement
- 2 Our Data
- 3 CNN Models
- 4 MobileNet Model
- 5 VGG Models
- 5 Model Performance
- 6 Conclusion & Future Work

Problem Statement

• • • •

As mask regulations continue to evolve it is more important than ever for governments and businesses to recognize individuals following and not following the appropriate safety guidelines. Based thereon, our goal was to create a solution to better enable the aforementioned stakeholders to identify which individuals are and are not wearing masks



Our Data

○ ○ ○ ○

Face Mask Detection Dataset

<https://www.kaggle.com/ashishjangra27/face-mask-12k-images-dataset>

Train:

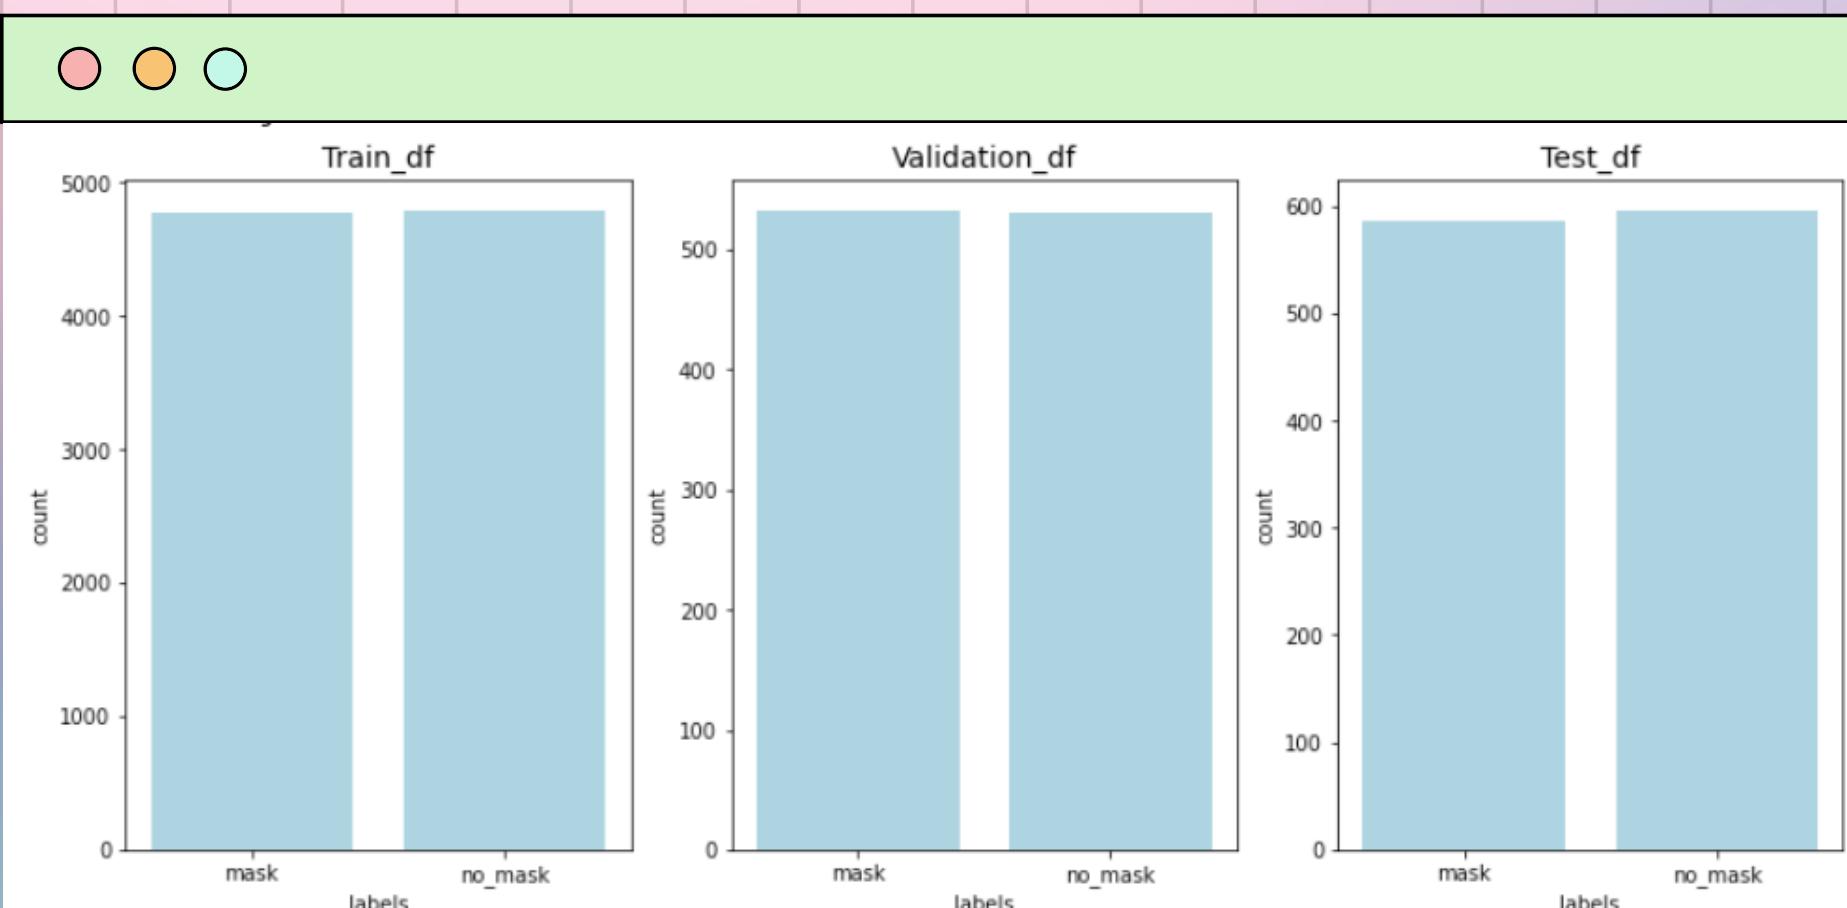
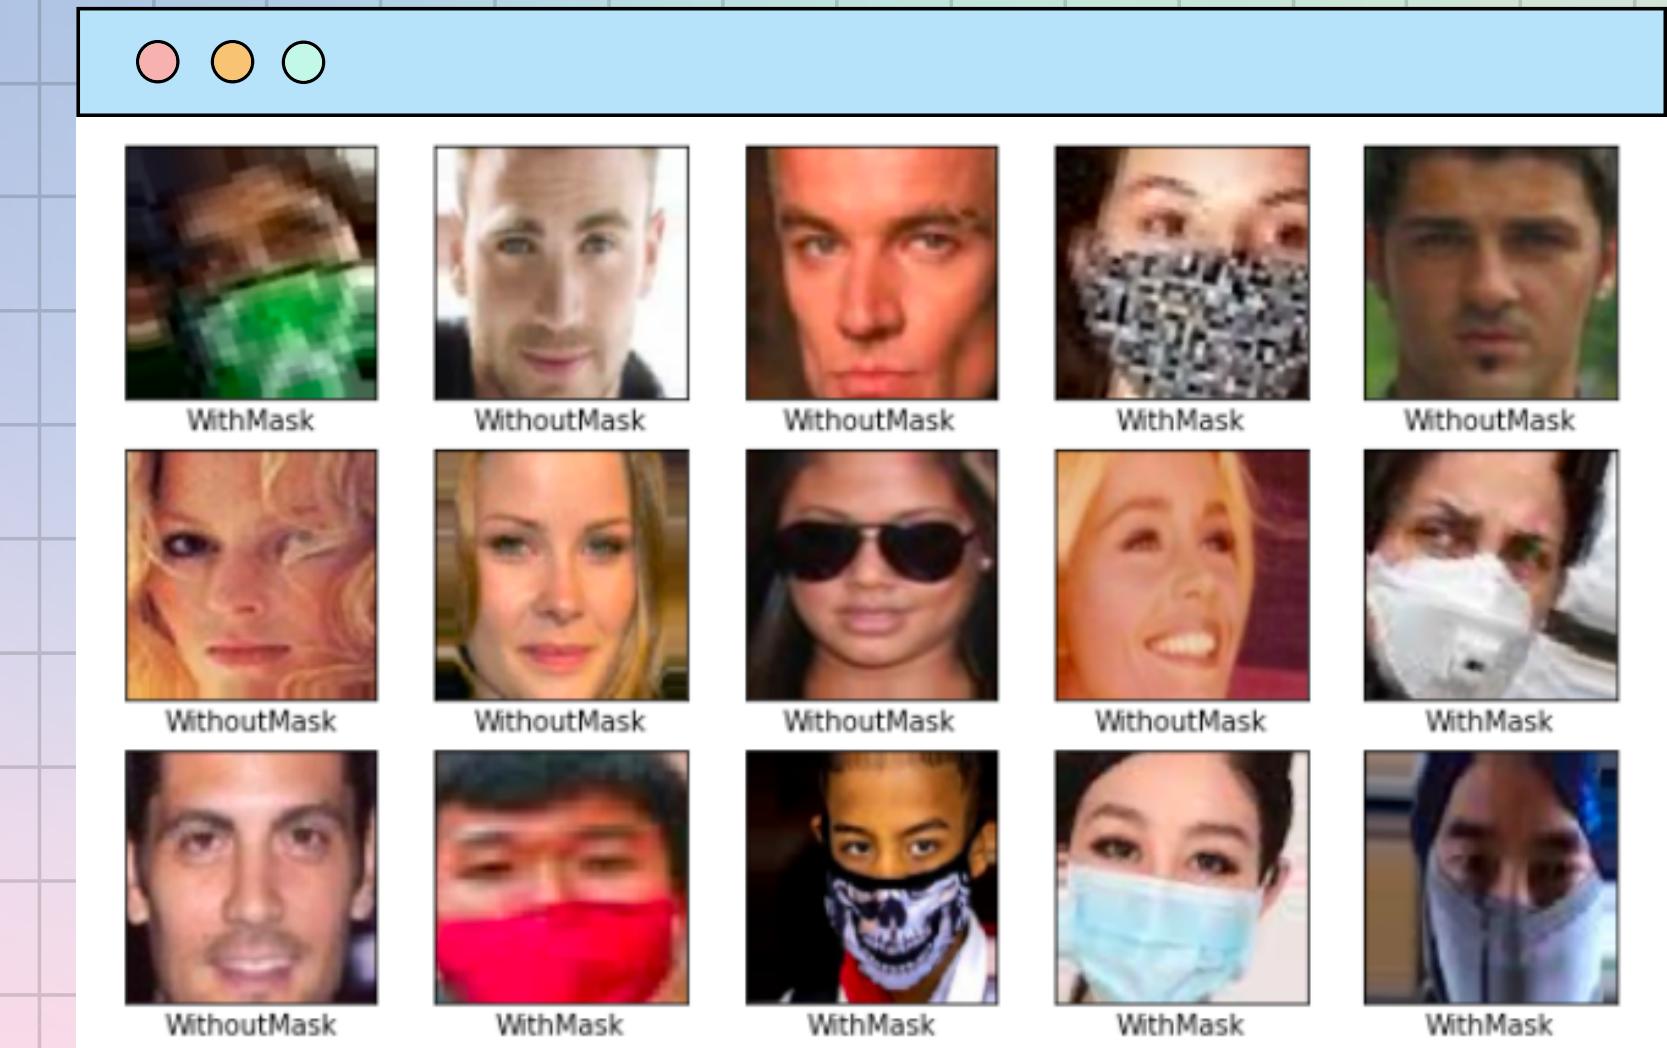
- No mask 4785
- Mask 4773

Validation:

- No mask 536
- Mask 526

Test:

- No_mask 597
- Mask 584



Haar Cascade

o o o o



haarcascade_frontalface_default.xml

Using haar cascade to detect faces

Object Detection using Haar feature-based cascade classifiers is an effective object detection method proposed by Paul Viola and Michael Jones in their paper, "Rapid Object Detection using a Boosted Cascade of Simple Features" in 2001.

- It is a machine learning based approach where a cascade function is trained from a lot of positive and negative images.
- Used to detect objects in other images.
- We used a Haar Cascade Model trained to detect faces in order to obtain the bounding box coordinates of faces in an image.

CNN Models

o o o o

Convolutional Neural Networks

o o o o



What is a CNN?

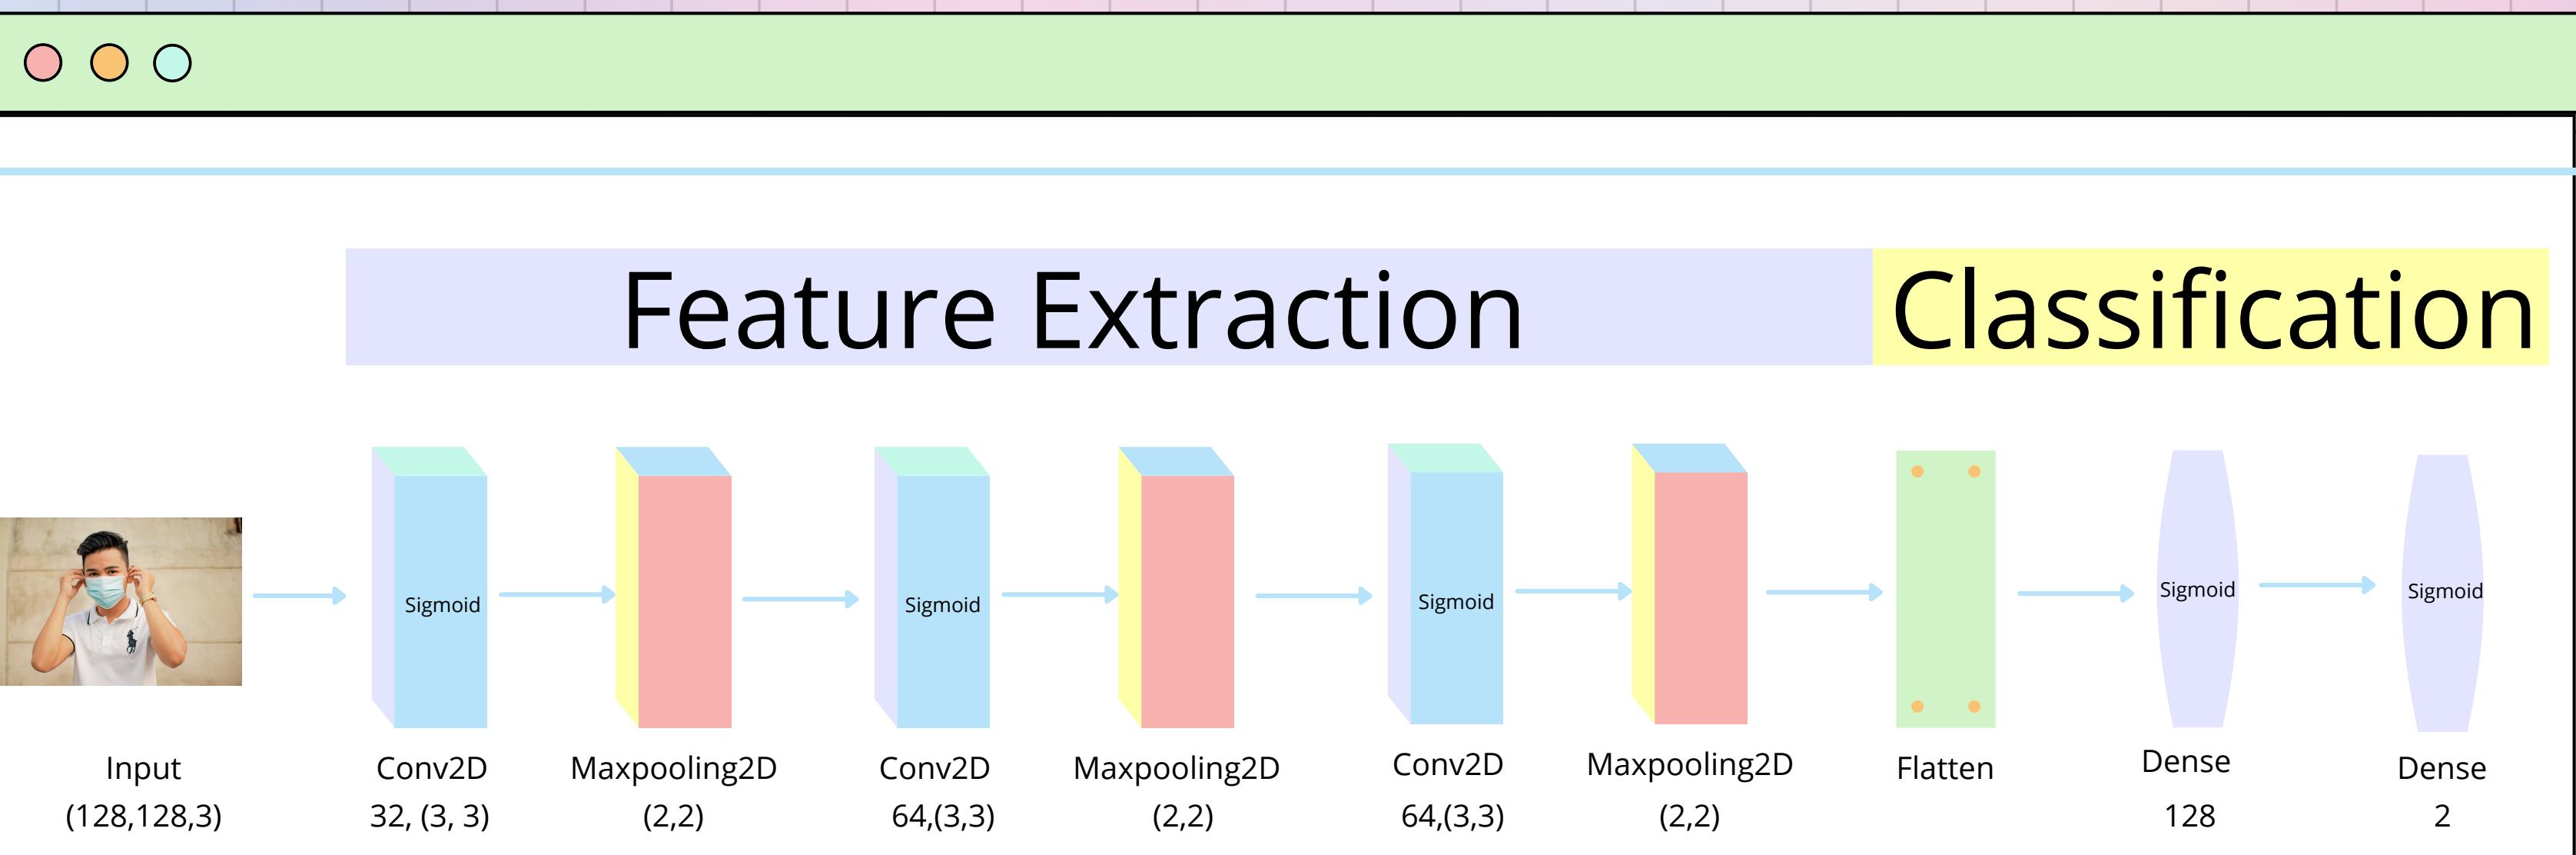


Assign importance
(learnable weights and
biases) to various aspects /
objects in the image

Classify

CNN Model 1:

○ ○ ○ ○

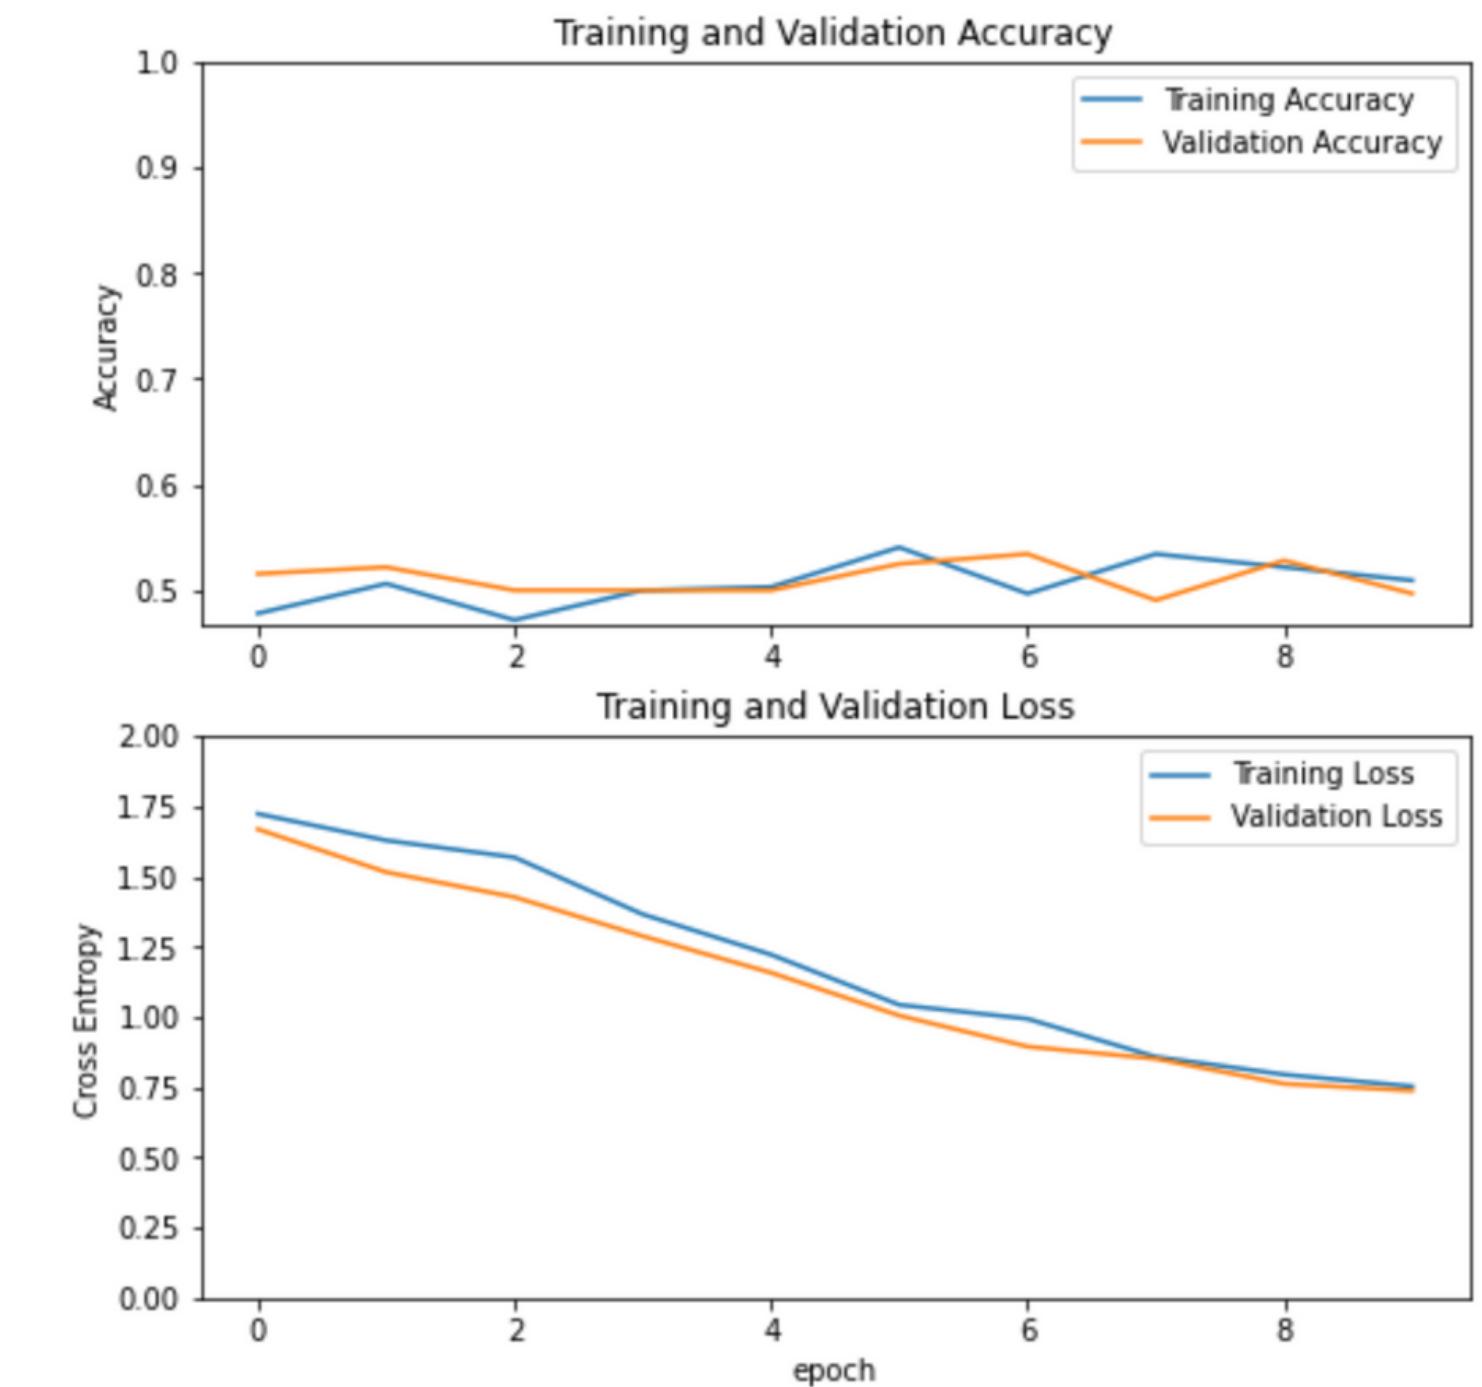


Performance

CNN Model 1

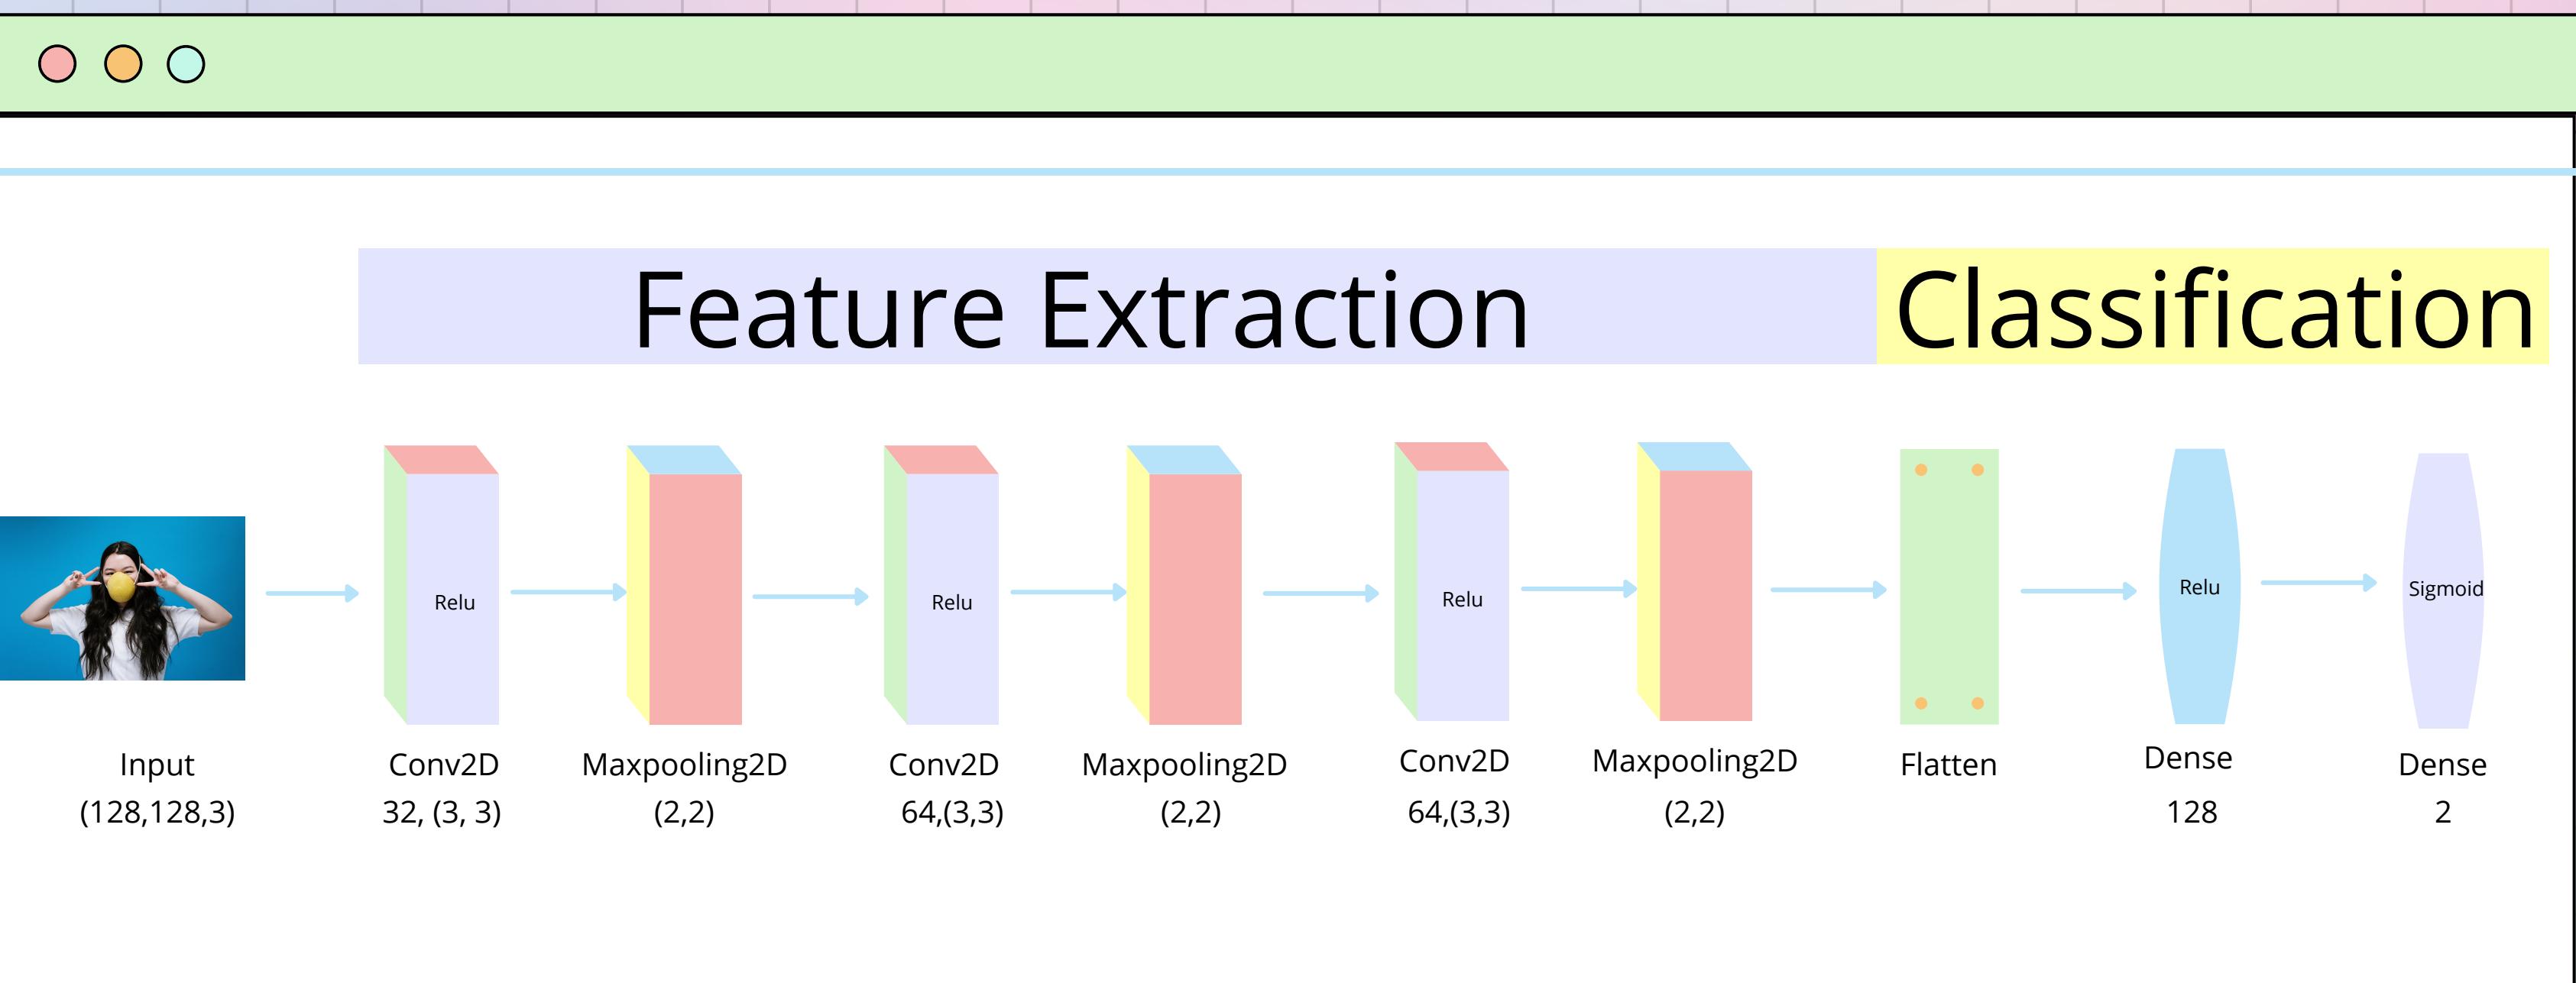
○ ○ ○ ○

Test Accuracy: **49.2%**



CNN Model 2:

o o o o



Performance

CNN Model 2

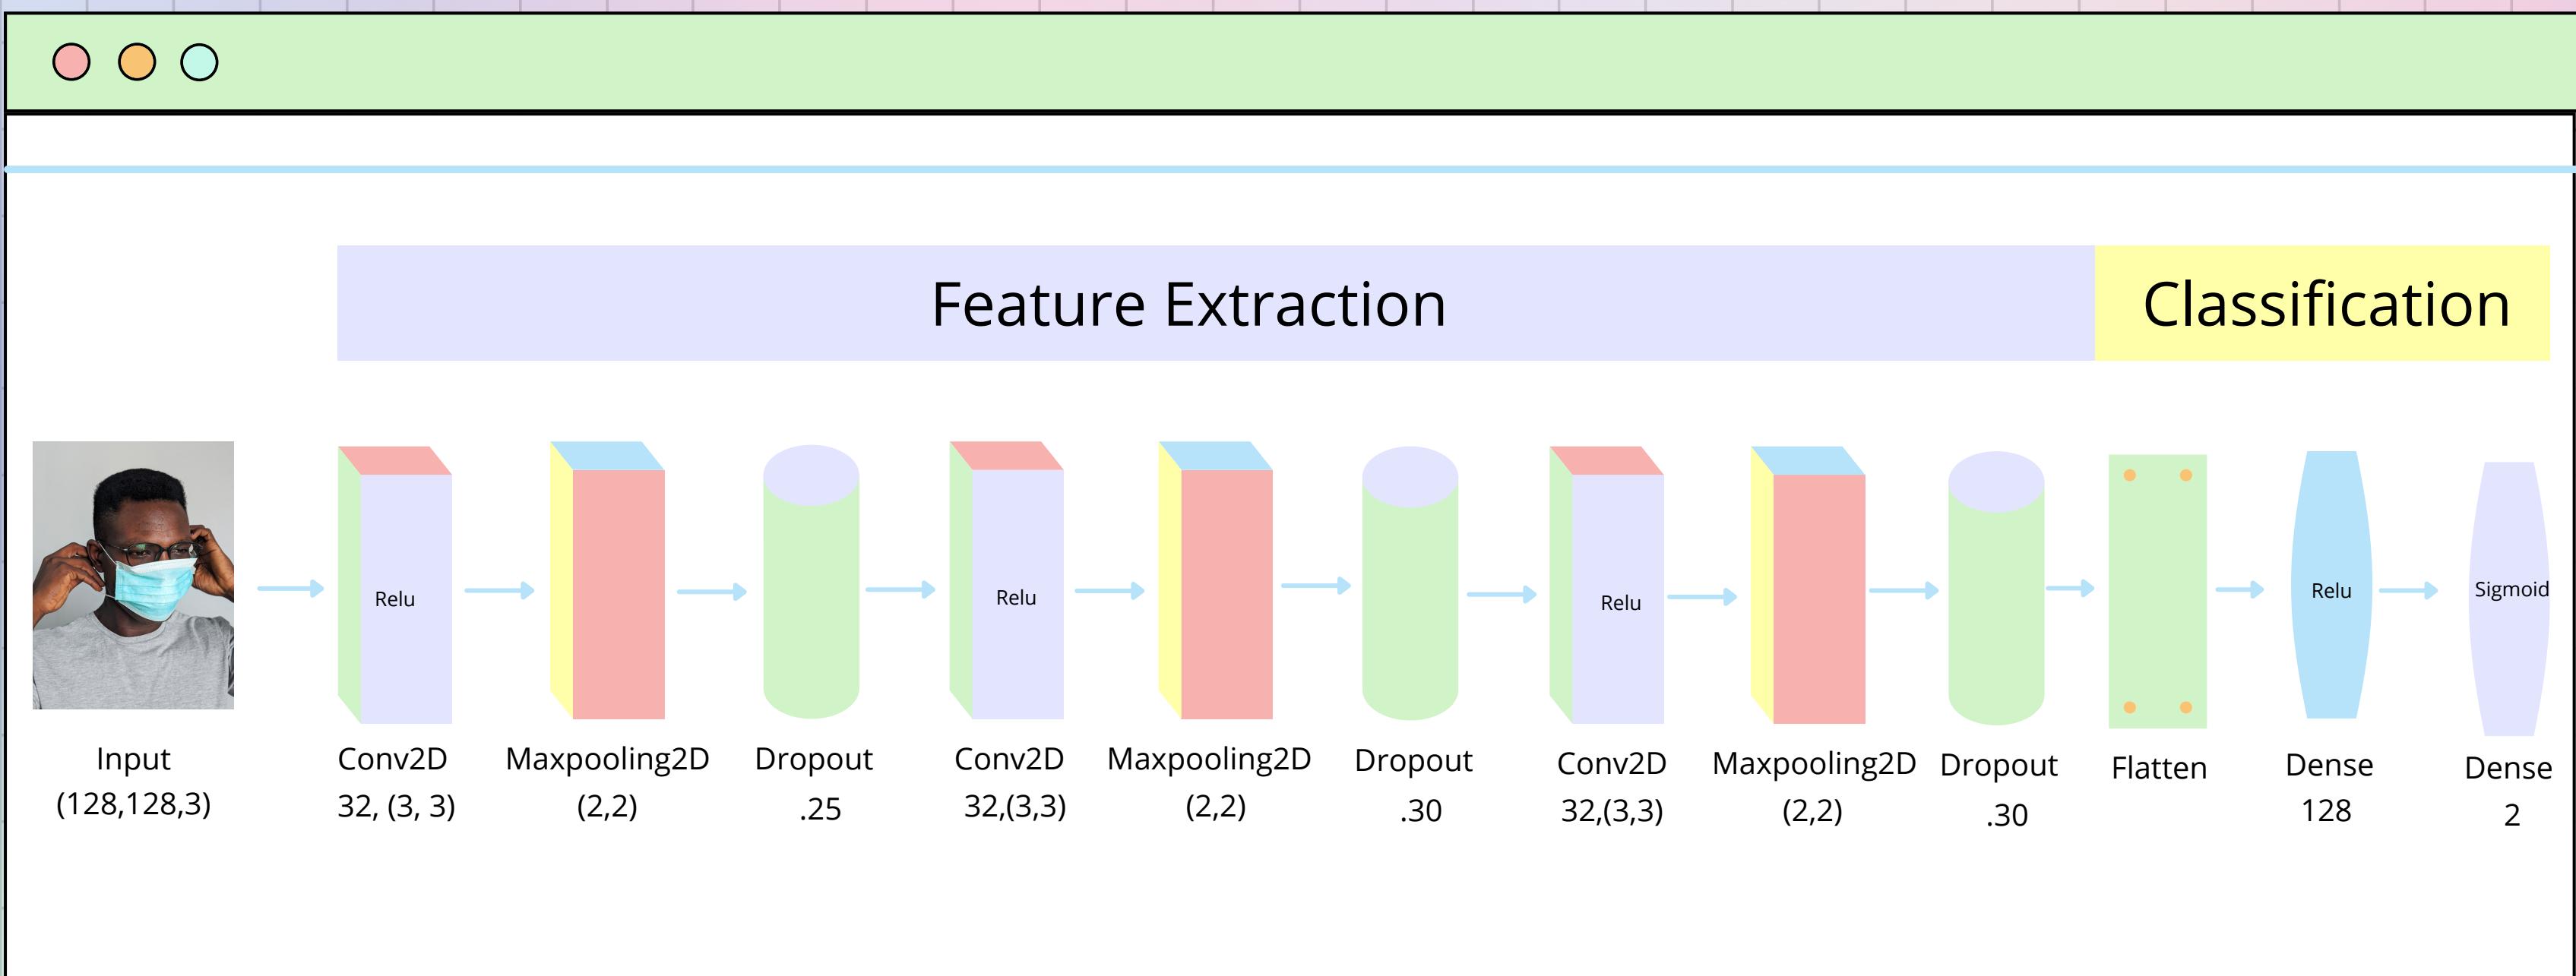
○ ○ ○ ○

Test Accuracy: **86.4%**



CNN Model 3:

o o o o



Performance

CNN Model 3

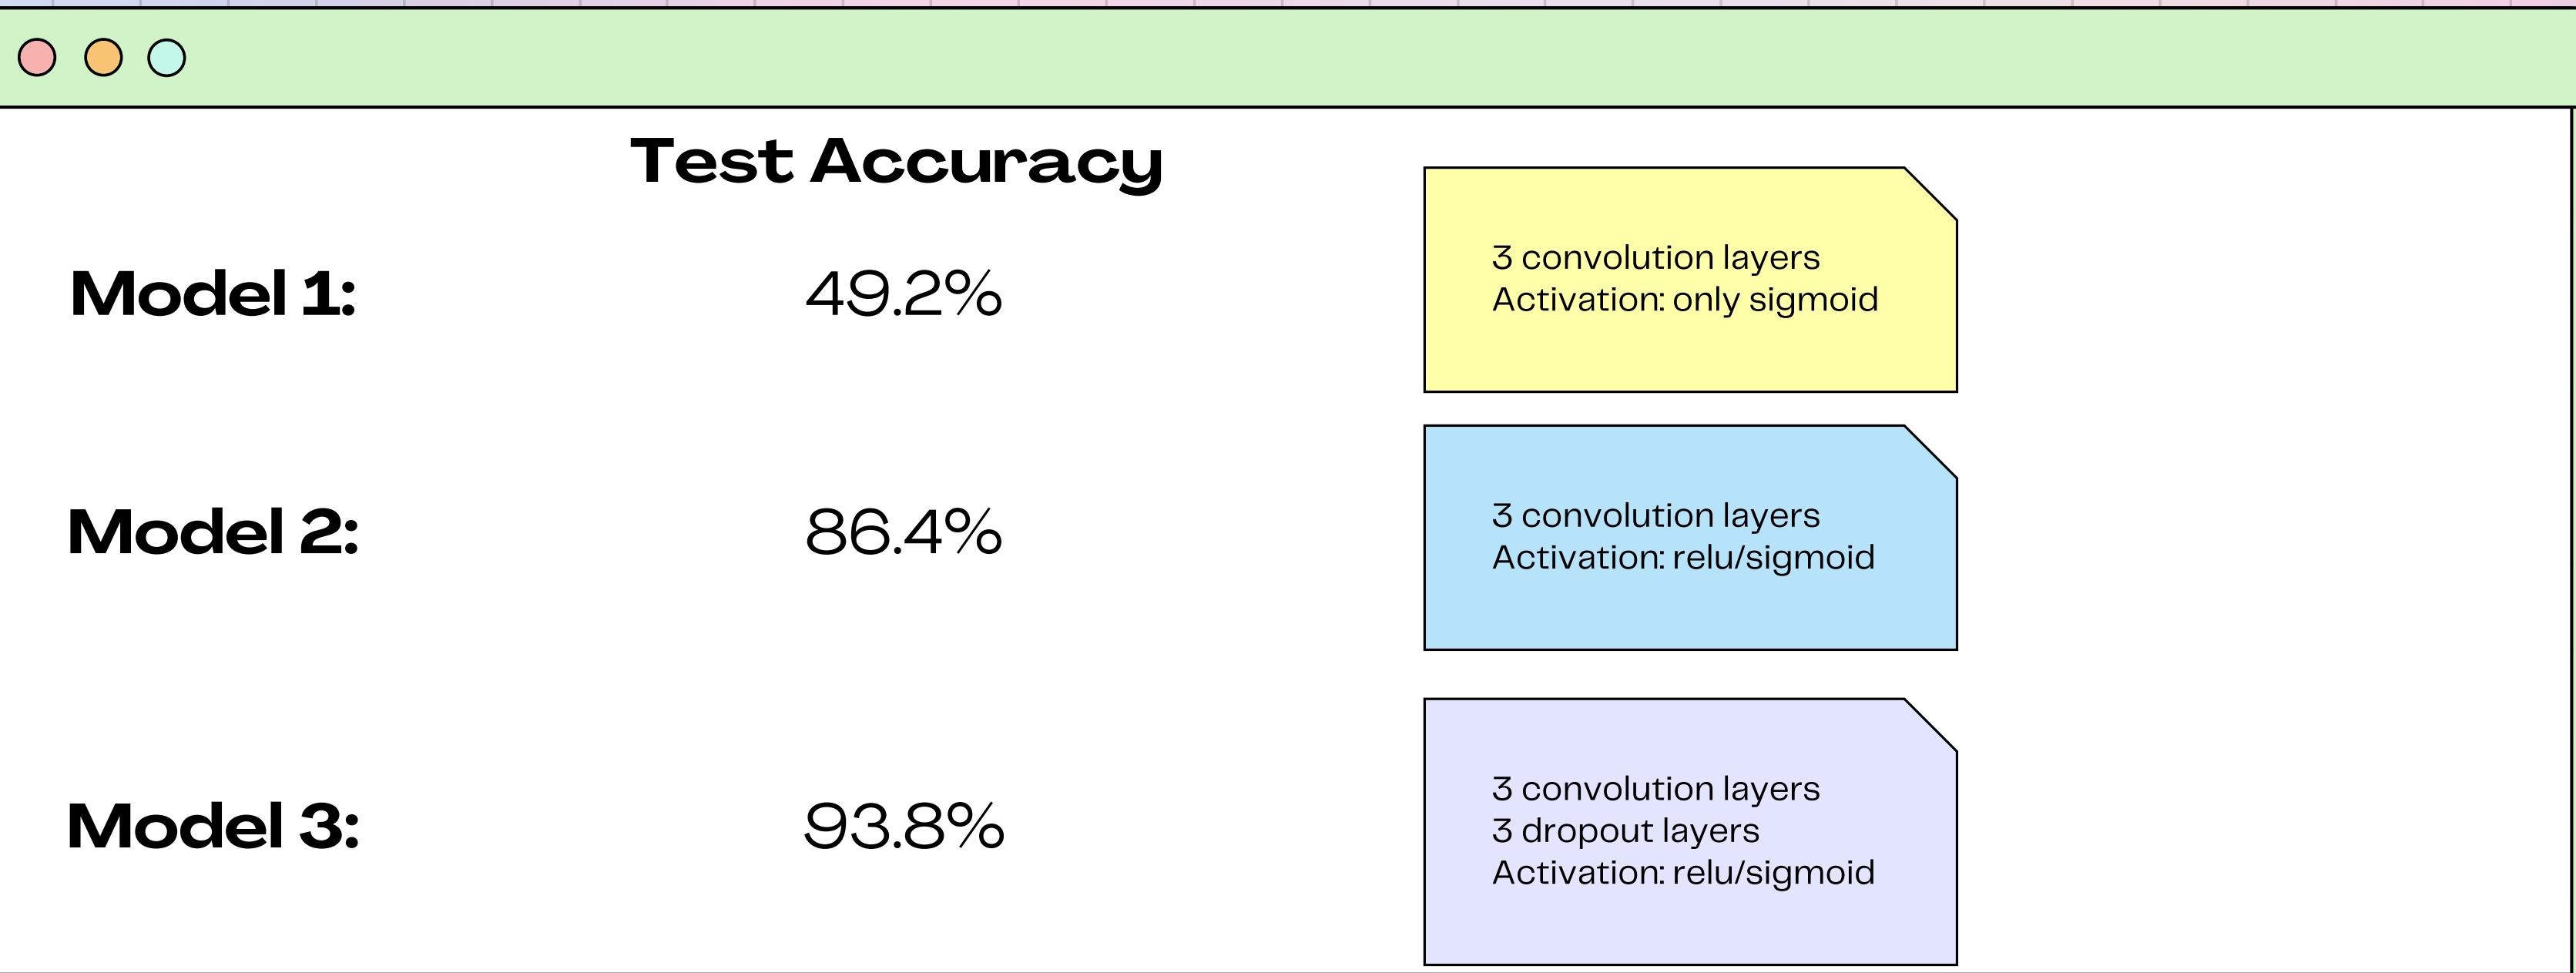
○ ○ ○ ○

Test Accuracy: **93.8%**



CNN Model Performance

○ ○ ○ ○

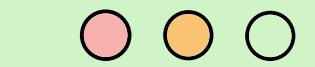
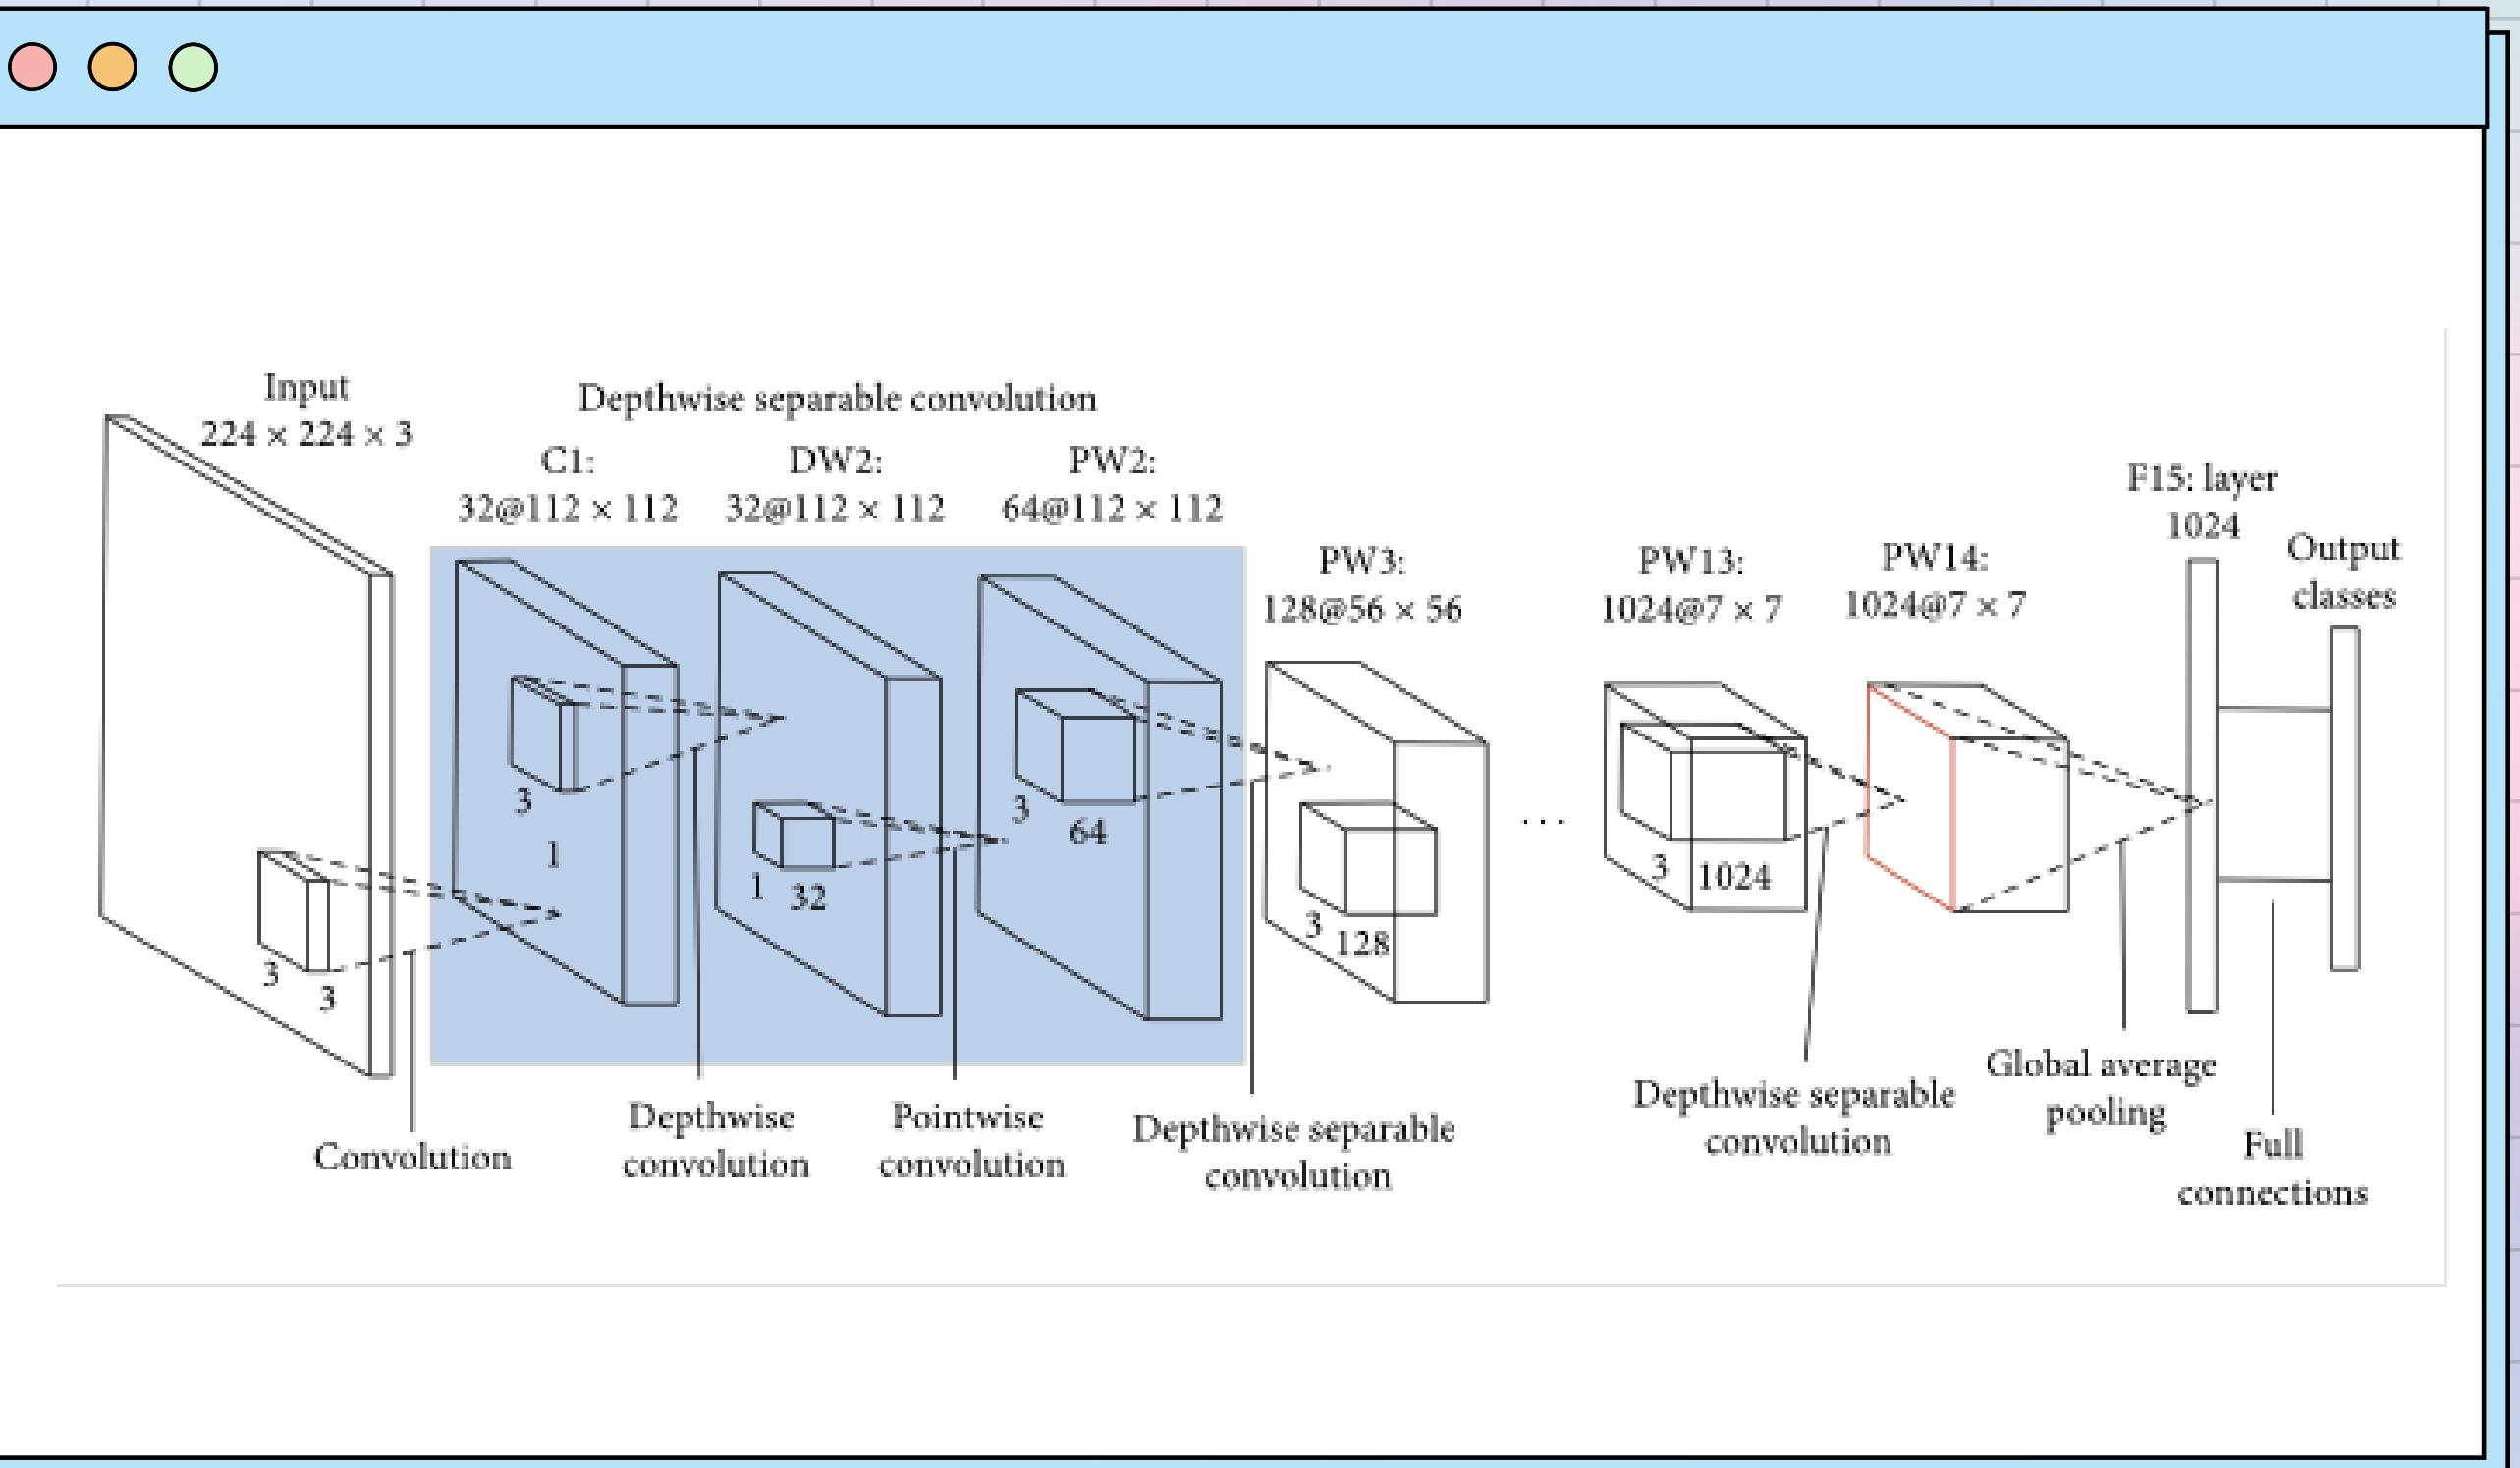


MobileNet Models

o o o o

MobileNet Model

○ ○ ○ ○



- depthwise separable convolutions
 - depthwise convolution
 - pointwise convolution
- Faster than the network with regular convolutions

MobileNet Model

o o o o



Type / Stride	Filter Shape
Conv / s2	$3 \times 3 \times 3 \times 32$
Conv dw / s1	$3 \times 3 \times 32$ dw
Conv / s1	$1 \times 1 \times 32 \times 64$
Conv dw / s2	$3 \times 3 \times 64$ dw
Conv / s1	$1 \times 1 \times 64 \times 128$
Conv dw / s1	$3 \times 3 \times 128$ dw
Conv / s1	$1 \times 1 \times 128 \times 128$
Conv dw / s2	$3 \times 3 \times 128$ dw
Conv / s1	$1 \times 1 \times 128 \times 256$
Conv dw / s1	$3 \times 3 \times 256$ dw
Conv / s1	$1 \times 1 \times 256 \times 256$
Conv dw / s2	$3 \times 3 \times 256$ dw
Conv / s1	$1 \times 1 \times 256 \times 512$
5× Conv dw / s1	$3 \times 3 \times 512$ dw
5× Conv / s1	$1 \times 1 \times 512 \times 512$
Conv dw / s2	$3 \times 3 \times 512$ dw
Conv / s1	$1 \times 1 \times 512 \times 1024$
Conv dw / s2	$3 \times 3 \times 1024$ dw
Conv / s1	$1 \times 1 \times 1024 \times 1024$
Avg Pool / s1	Pool 7×7
FC / s1	1024×1000
Softmax / s1	Classifier

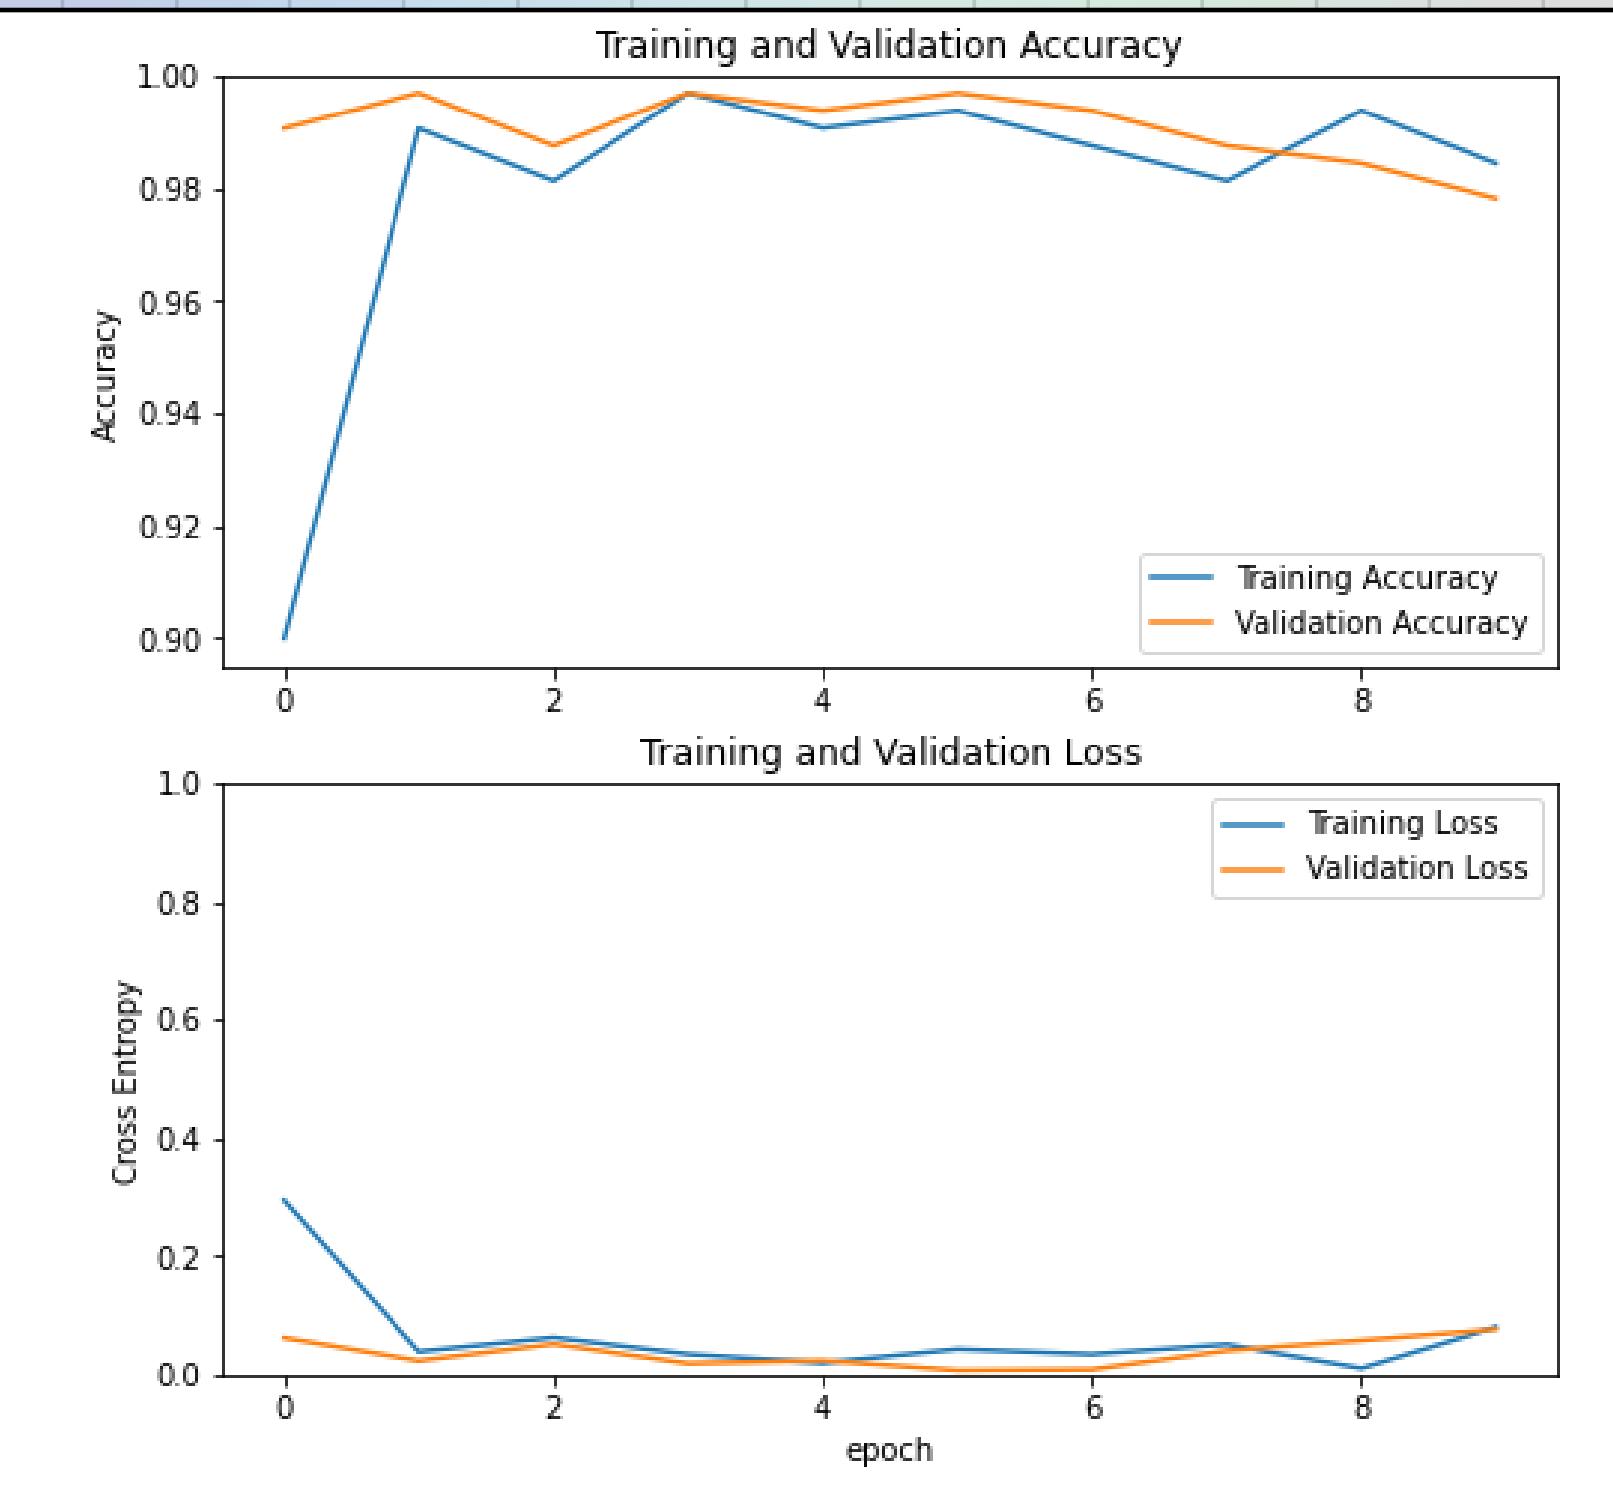
- Depthwise Separable Convolutions
- Convolution layer
- Average Pooling layer
- Fully Connected layer
- Total params: 3,261,634
- Trainable params: 32,770
- Non-trainable params: 3,228,864

Performance

MobileNet

○ ○ ○ ○

Test Accuracy: **98.6%**



VGG Models

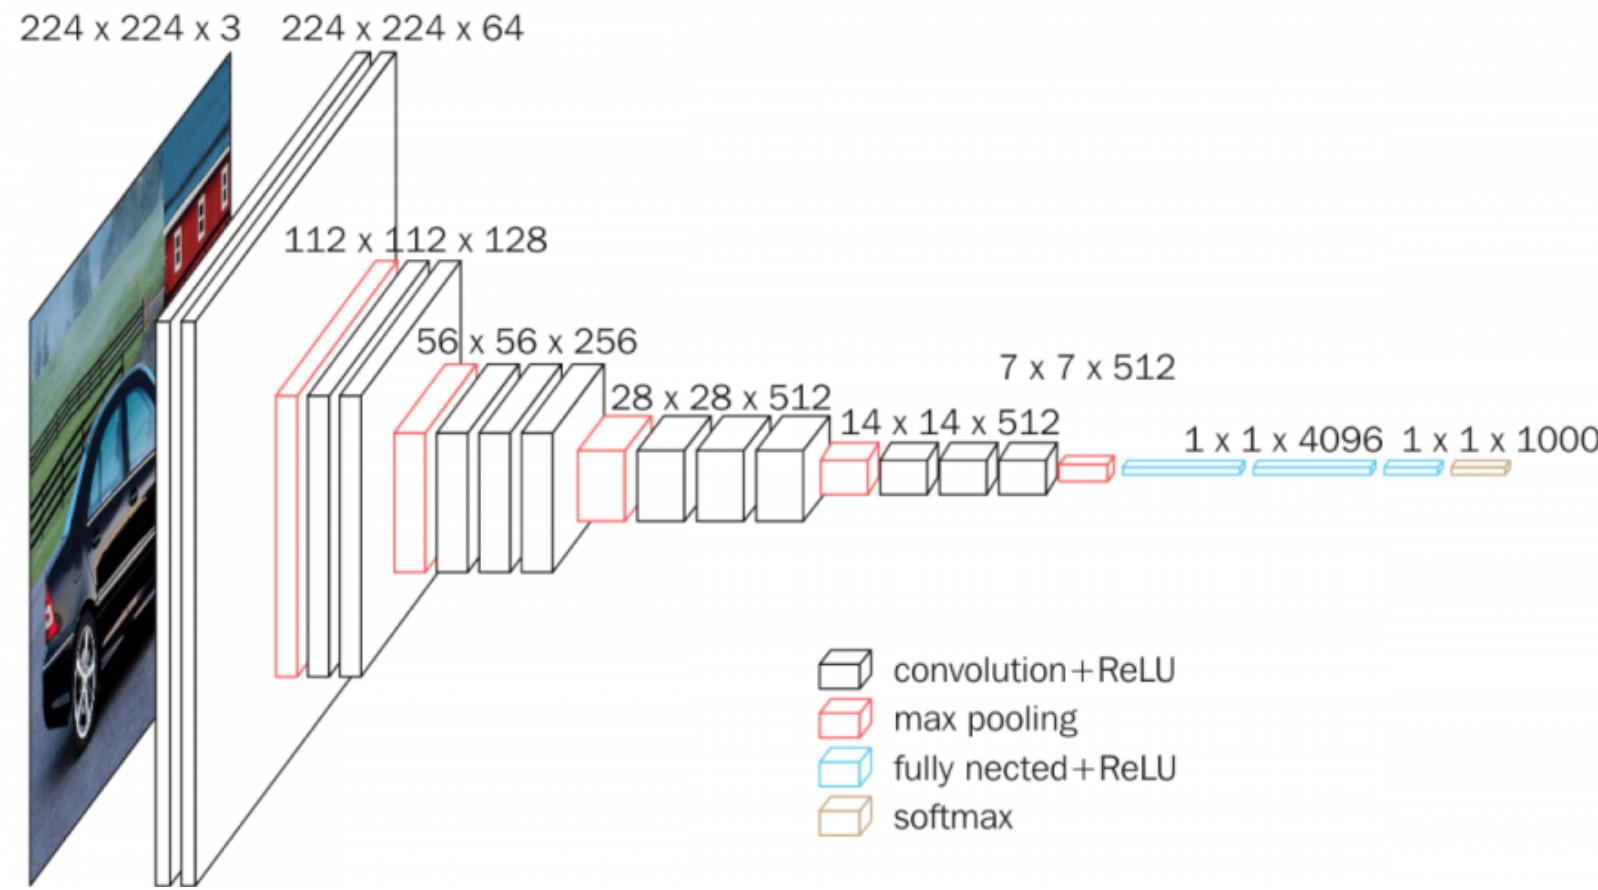
○ ○ ○ ○

Transfer Learning & Fine-tuning

Base Model

o o o o

Visual Geometry Group (VGG)



1

Trained on ImageNet dataset

2

Improvements based on AlexNet from large-sized kernels to multiple 3x3 kernels

Over 14 million images & 1000 classes

Fixed input size 224x224 RGB image

Painfully slow to train & Large Architecture

VGG 16

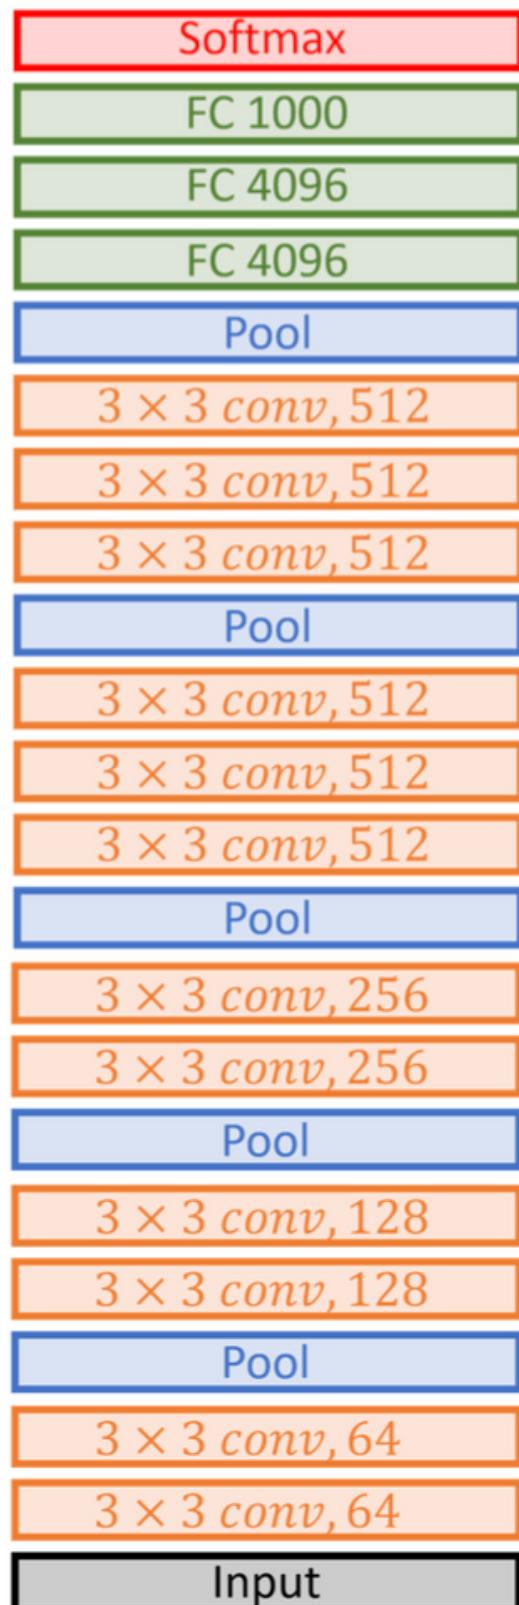
○ ○ ○ ○



1 16 weighted layers

2 Relatively Large Network Architecture

3 ReLU non-linearity



CONVOLUTIONAL LAYERS X12

- 3x3 Filter Size
- Stride 1

MAXPOOL LAYERS X5

- 2x2 Filter Size
- Stride 2

FULLY CONNECTED LAYERS X3

SOFTMAX LAYER X1

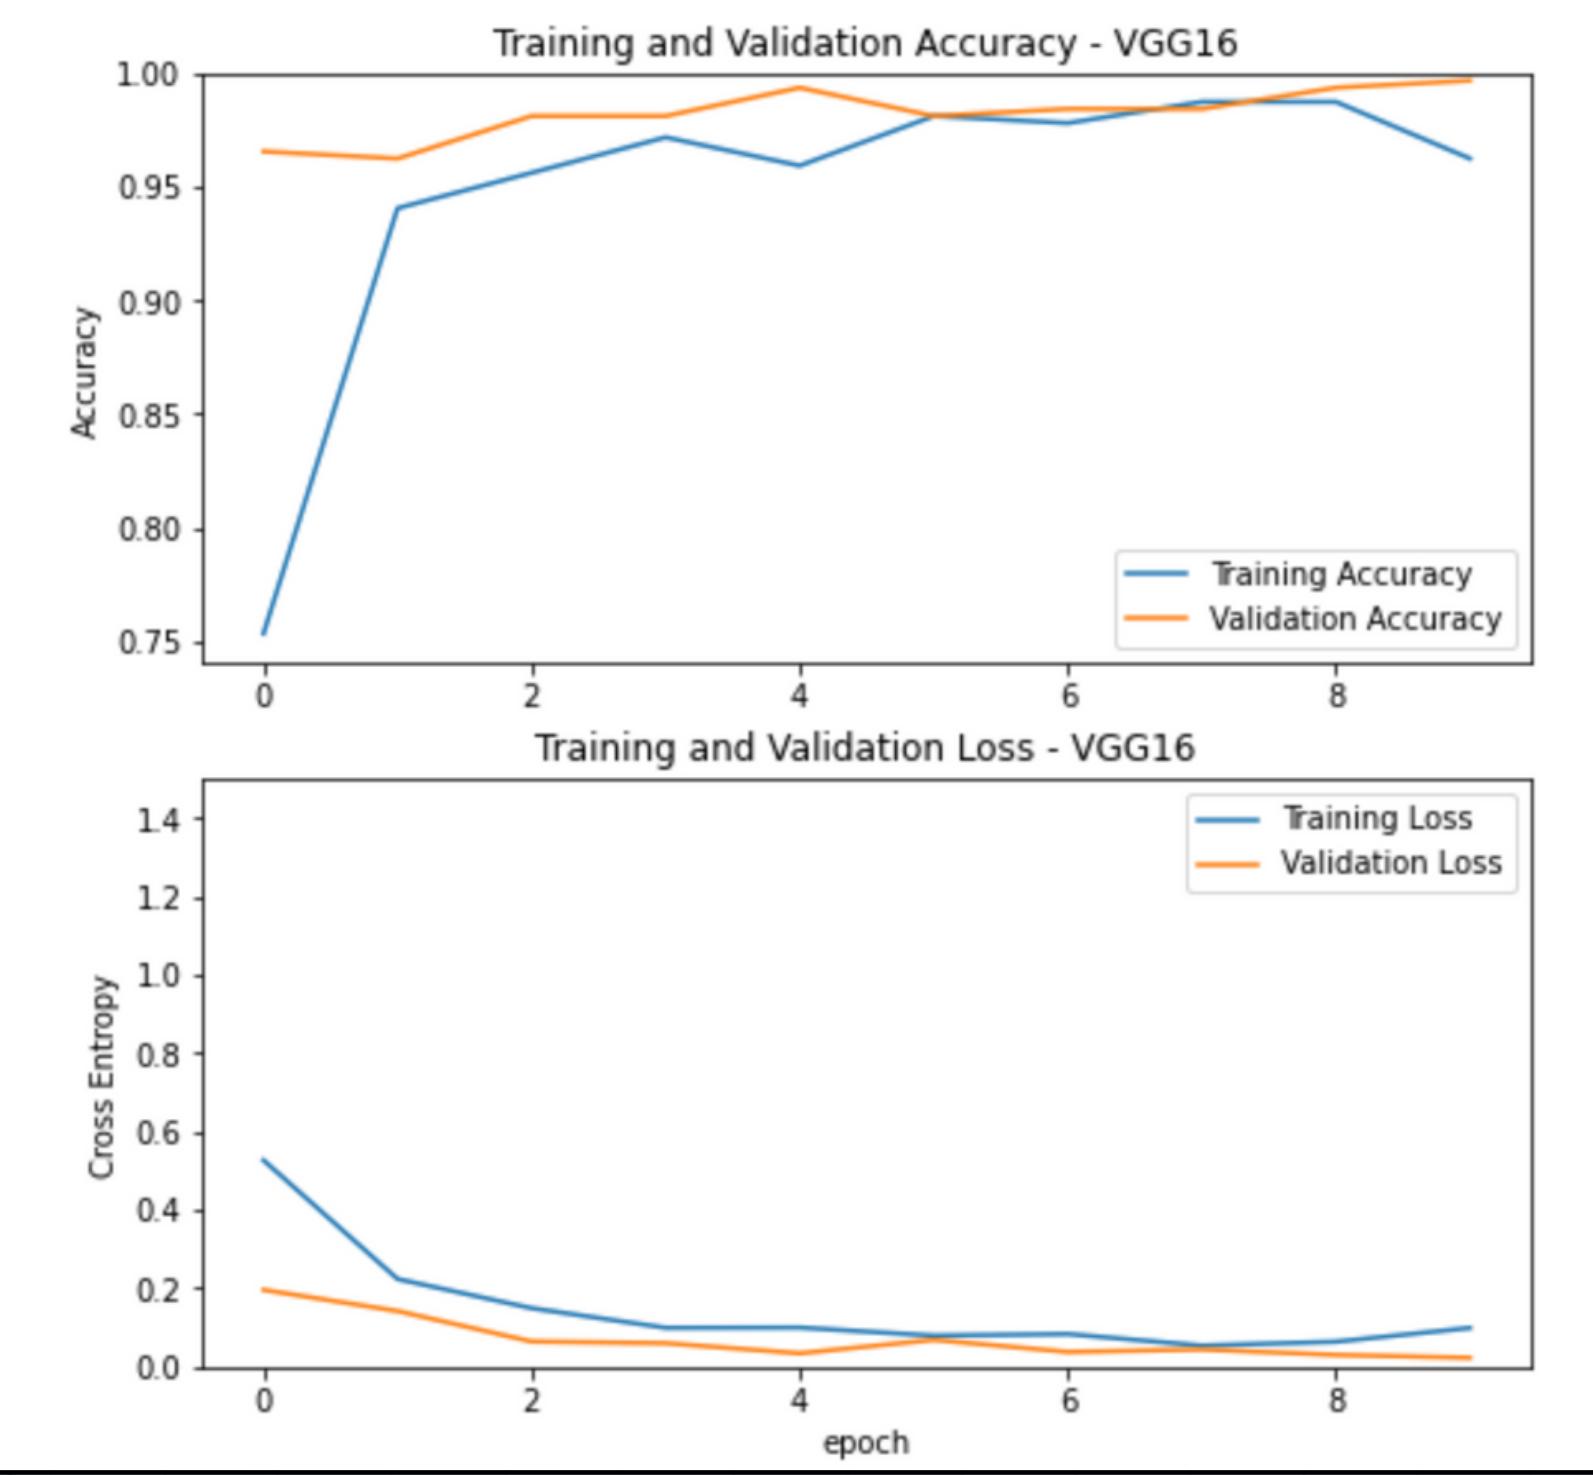
- Total parameters: 14,731,074
- Trainable parameters: 16,386
- Non-trainable parameters: 14,714,688

Performance

VGG16 Base Model

○ ○ ○ ○

Test Accuracy: 99.2%



VGG 19

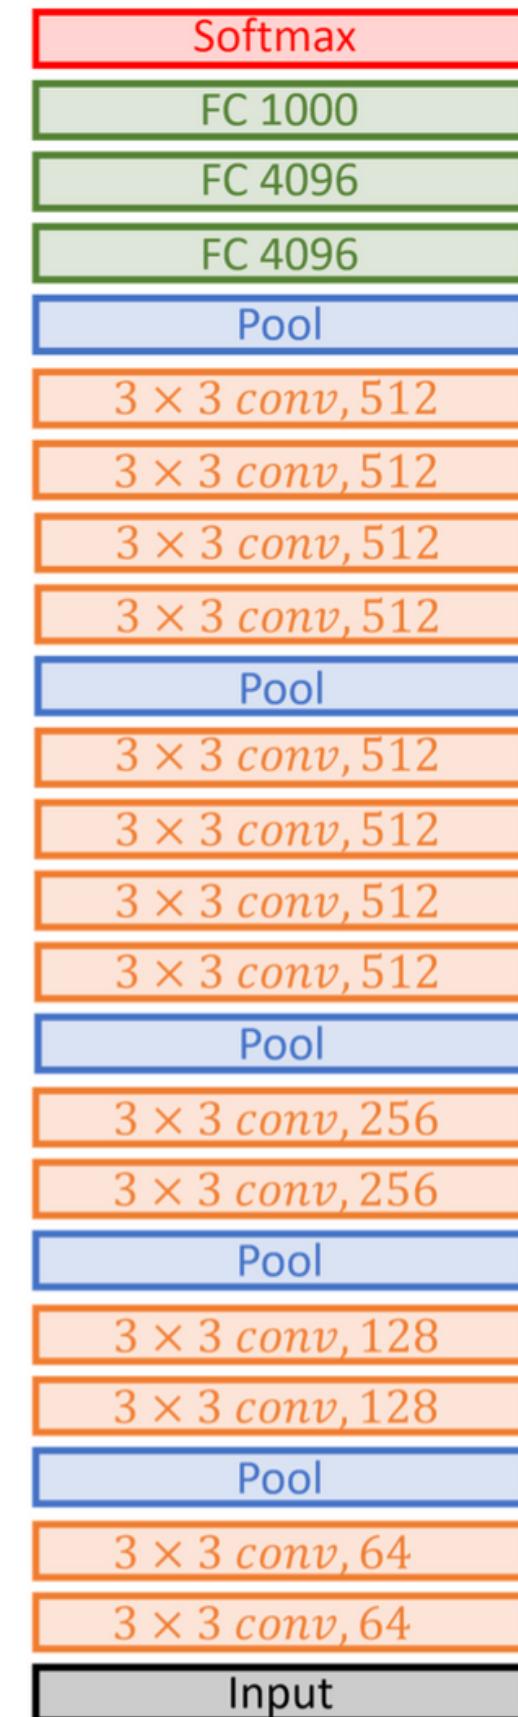
○ ○ ○ ○



1 19 Weighted Layers

2 +2 Convolutional Layers

3 Compare with VGG 16



CONVOLUTIONAL LAYERS X14

- 3x3 Filter Size
- Stride 1

MAXPOOL LAYERS X5

- 2x2 Filter Size
- Stride 2

FULLY CONNECTED LAYERS X3

SOFTMAX LAYER X1

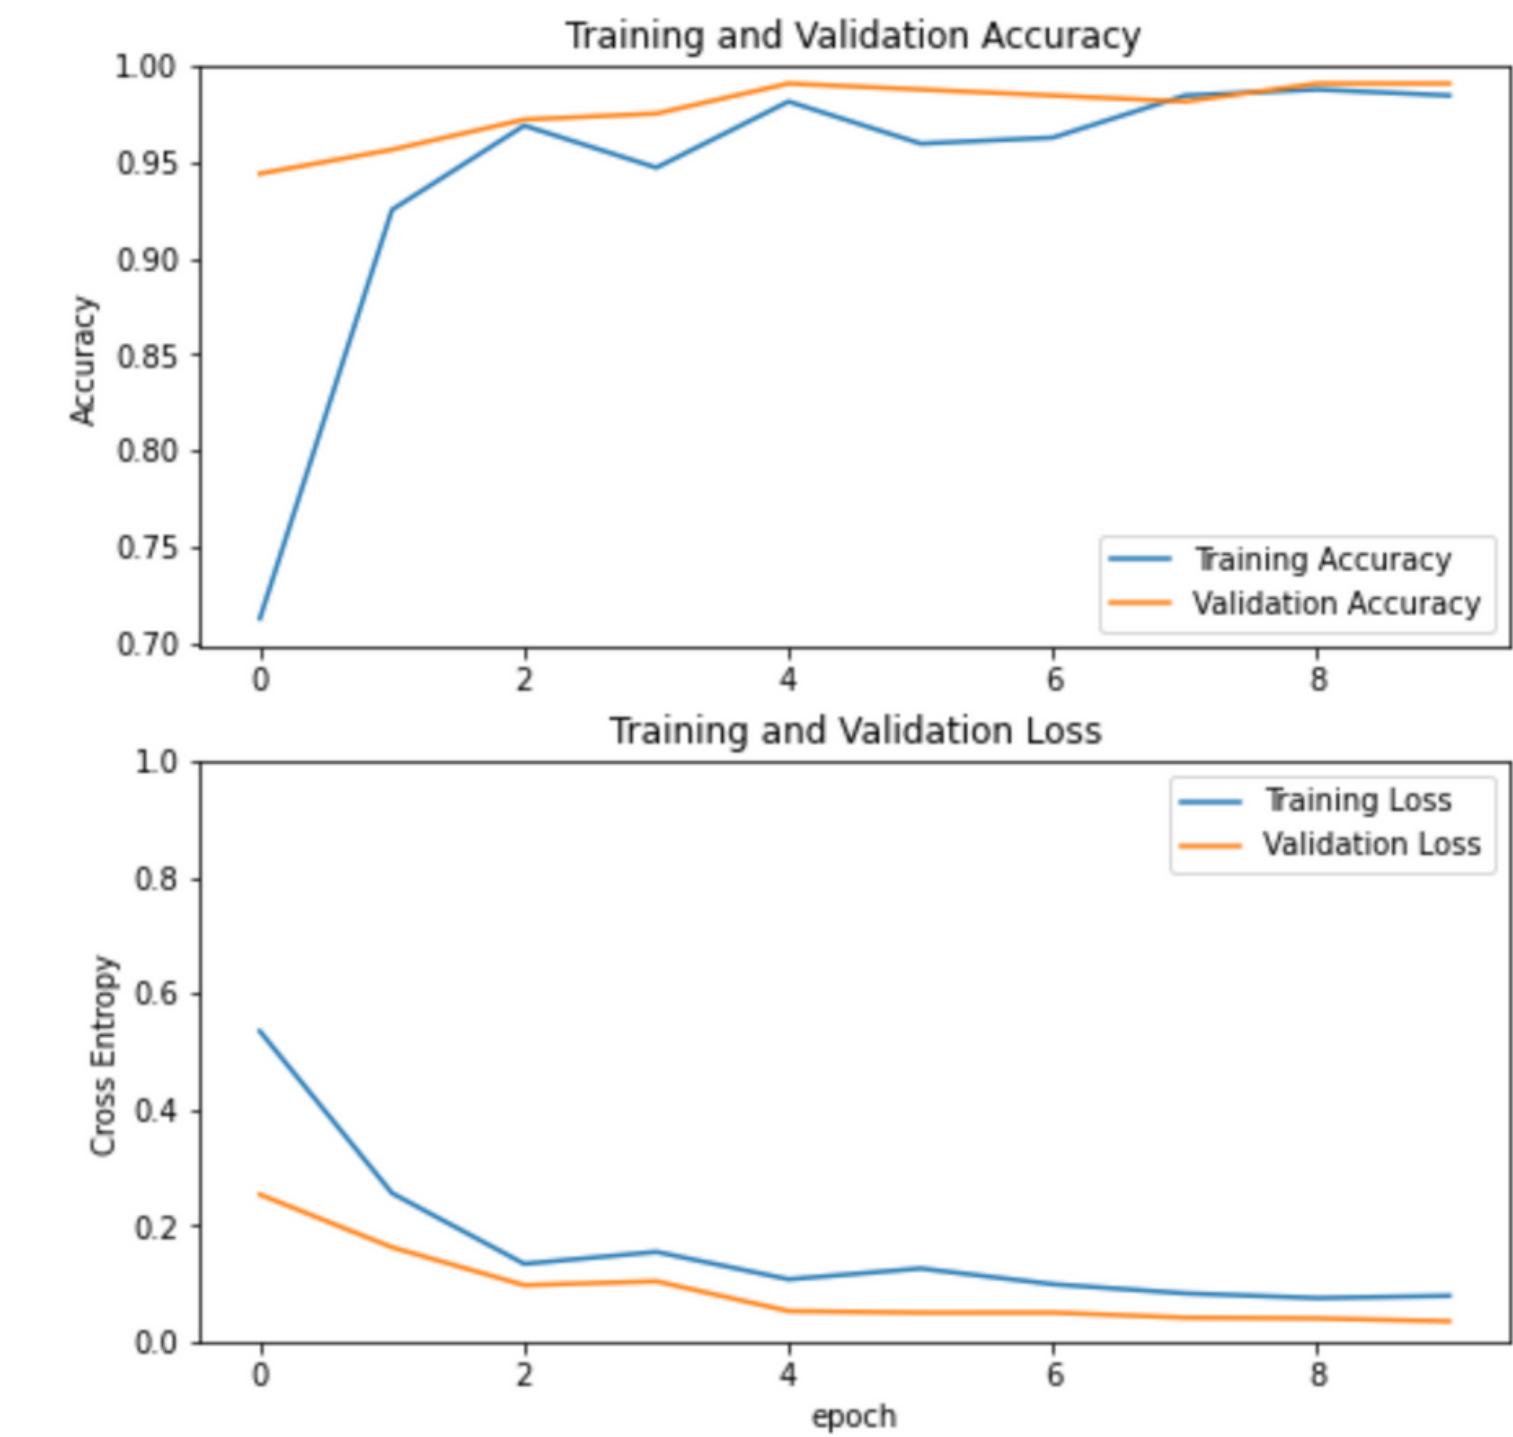
- Total parameters: 20,040,770
- Trainable parameters: 16,386
- Non-trainable parameters: 20,024,384

Performance

VGG19 Base Model

○ ○ ○ ○

Test Accuracy: 98.7%



Fine-tuning

○ ○ ○ ○



Action 1

Take layers from a previously trained model.



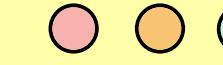
Action 2

Freeze them, so as to avoid destroying any of the information they contain during future training rounds.



Action 3

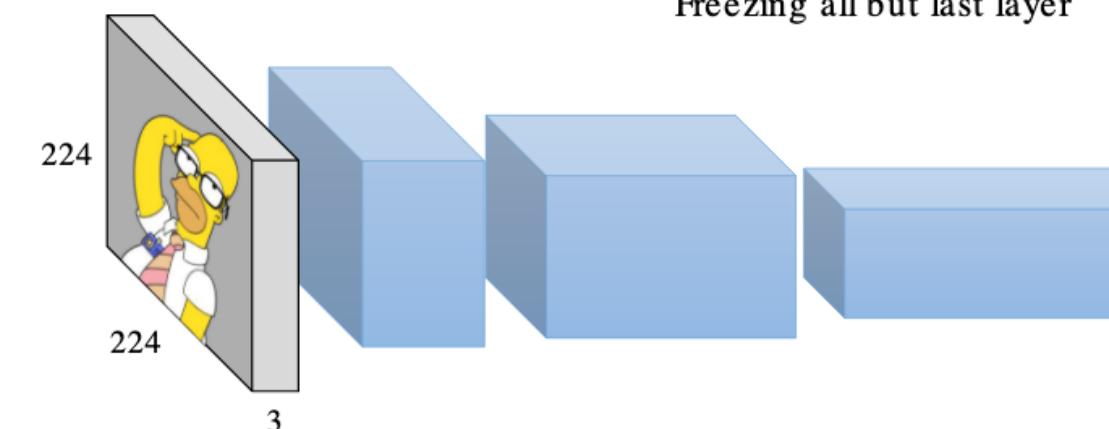
Add some new, trainable layers on top of the frozen layers. Train the new layers on your dataset.



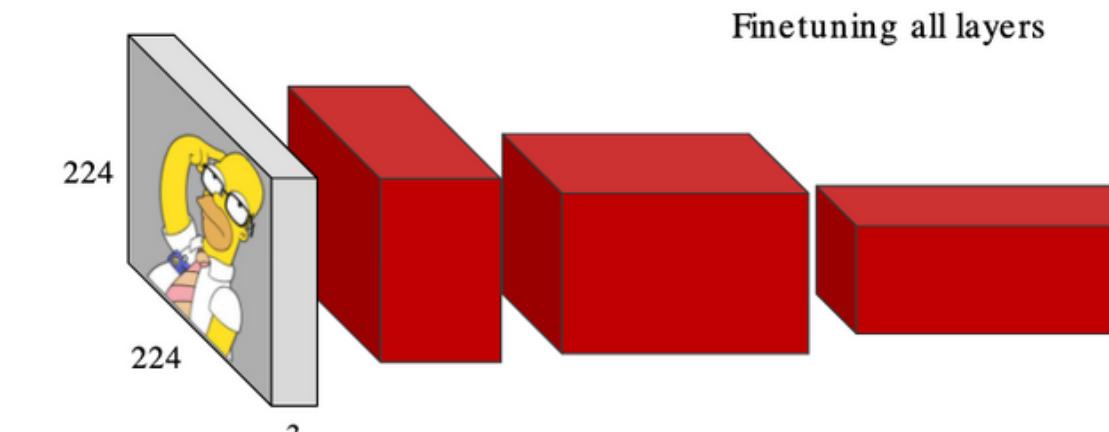
Action 4

Fine-tuning: consists of unfreezing the entire model you obtained above , and re-training it on the new data with a very low learning rate.

Freezing all but last layer



last layer
...
20 units



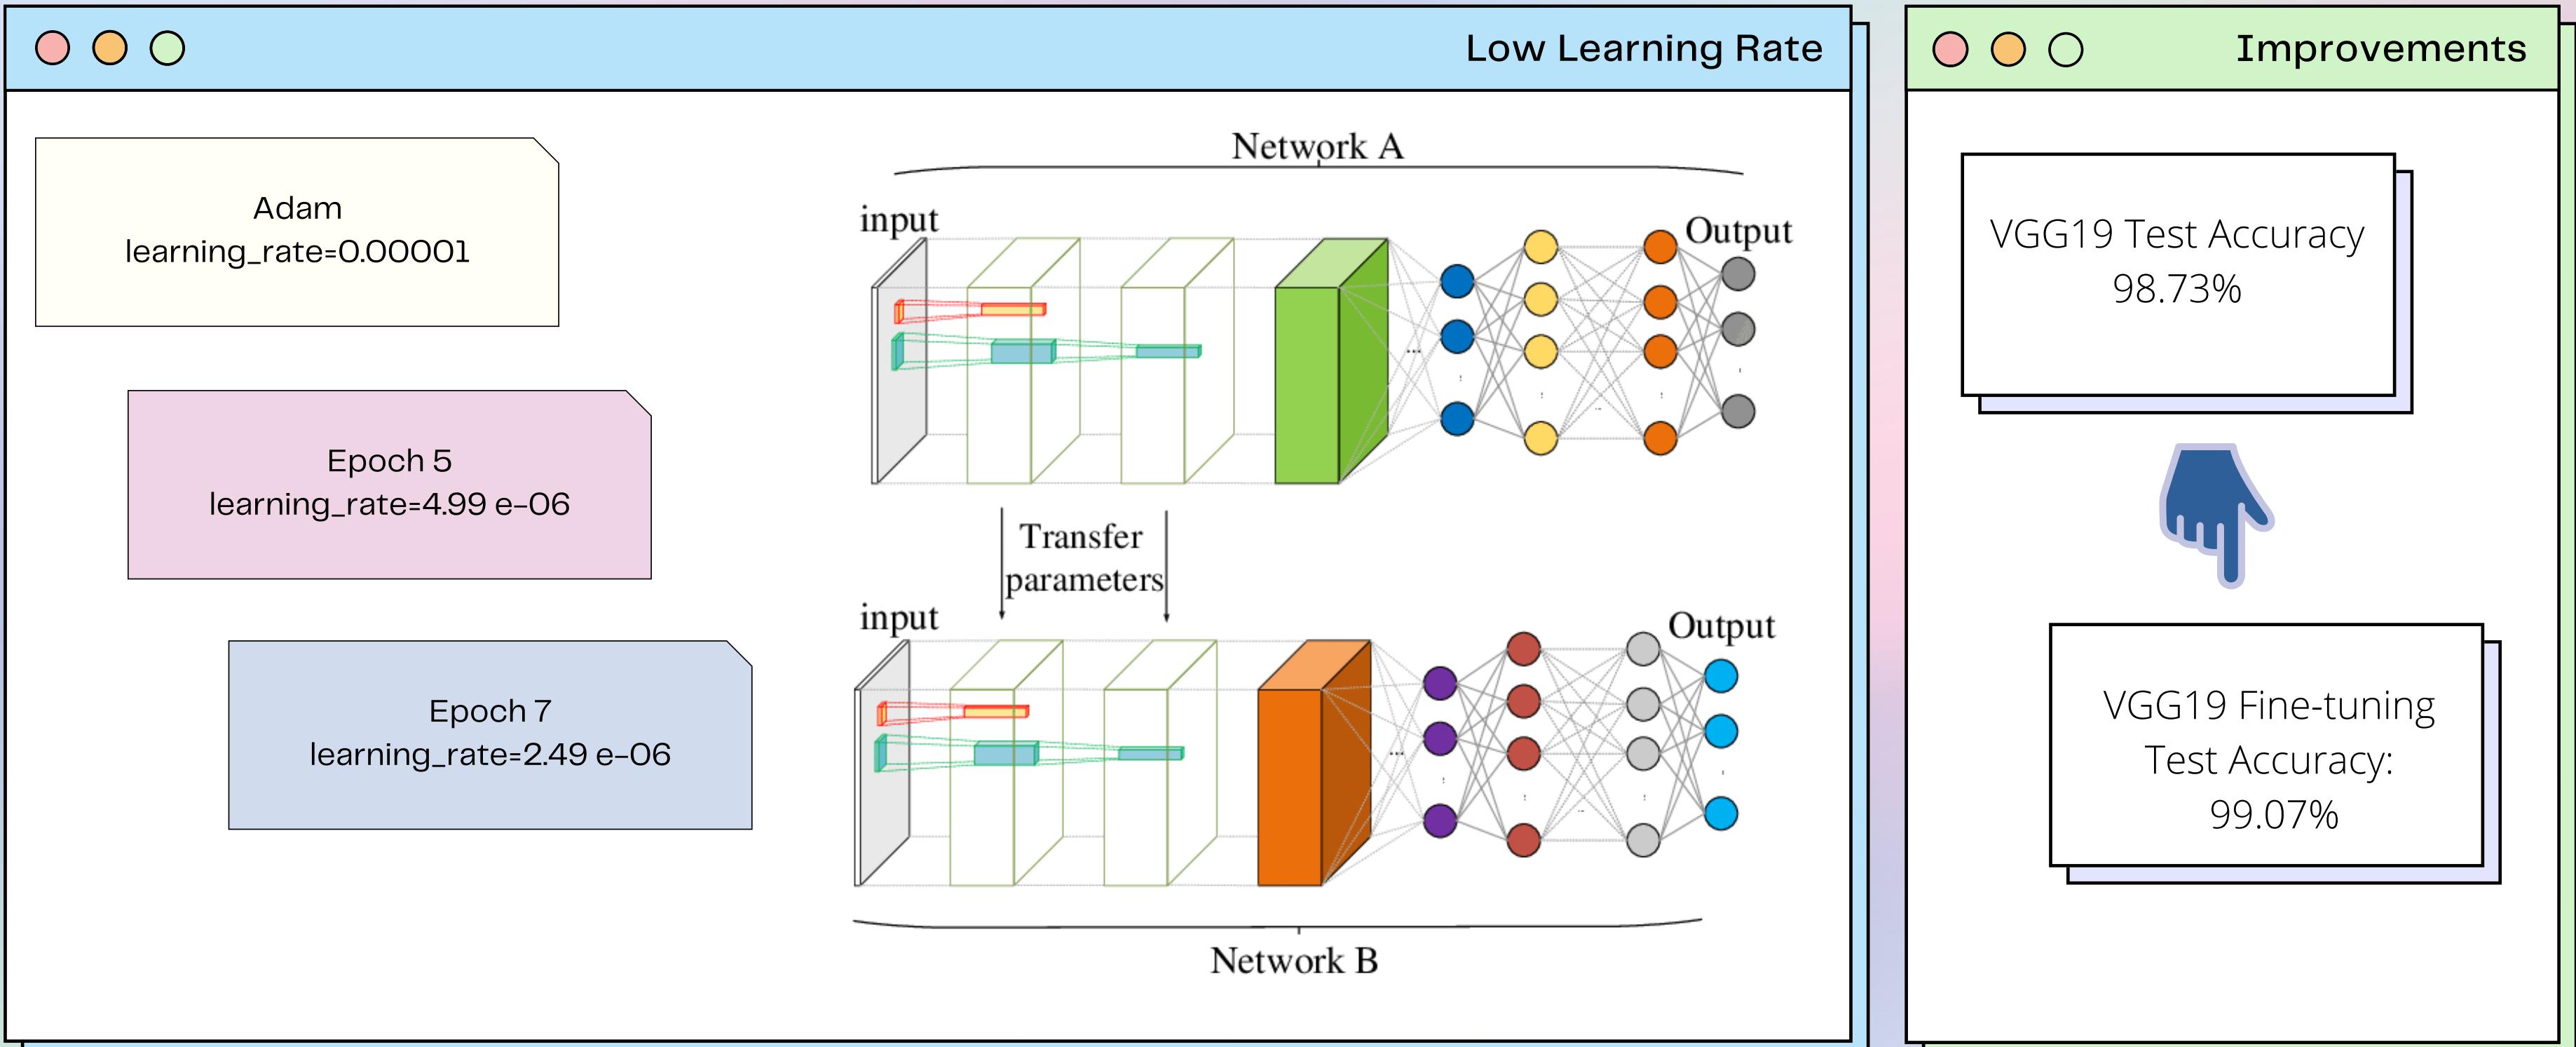
last layer
...
20 units

Fine-tuning & Callbacks

1 Fine-Tuning
Unfreeze all the layer & Update weights.

2 EarlyStopping
End training when val_accuracy is not improved for 4 consecutive times.

3 ReduceLROnPlateau
The learning rate is reduced by half when val_accuracy is not improved for 2 consecutive times

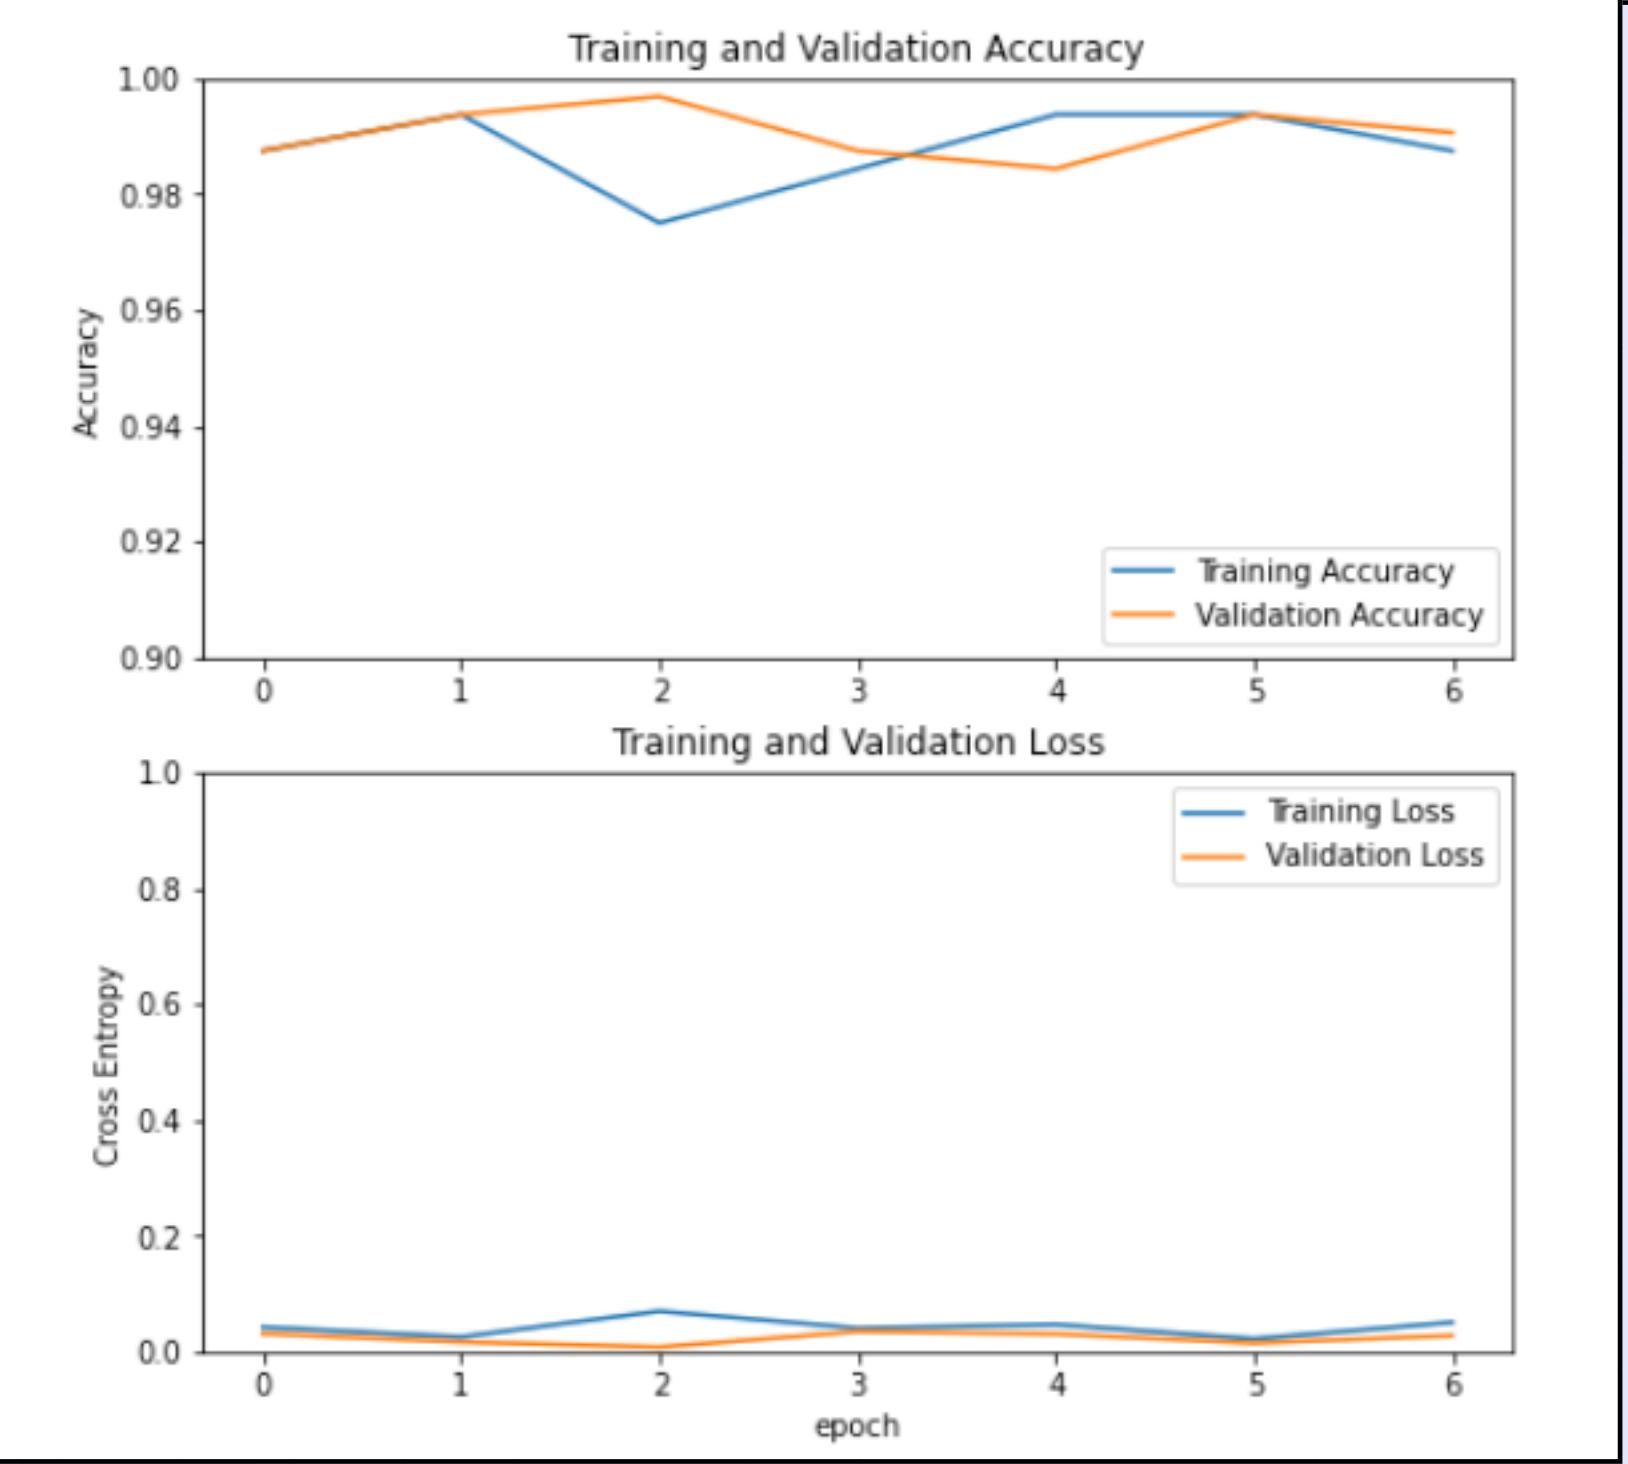


Performance

VGG19 Fine-Tuning

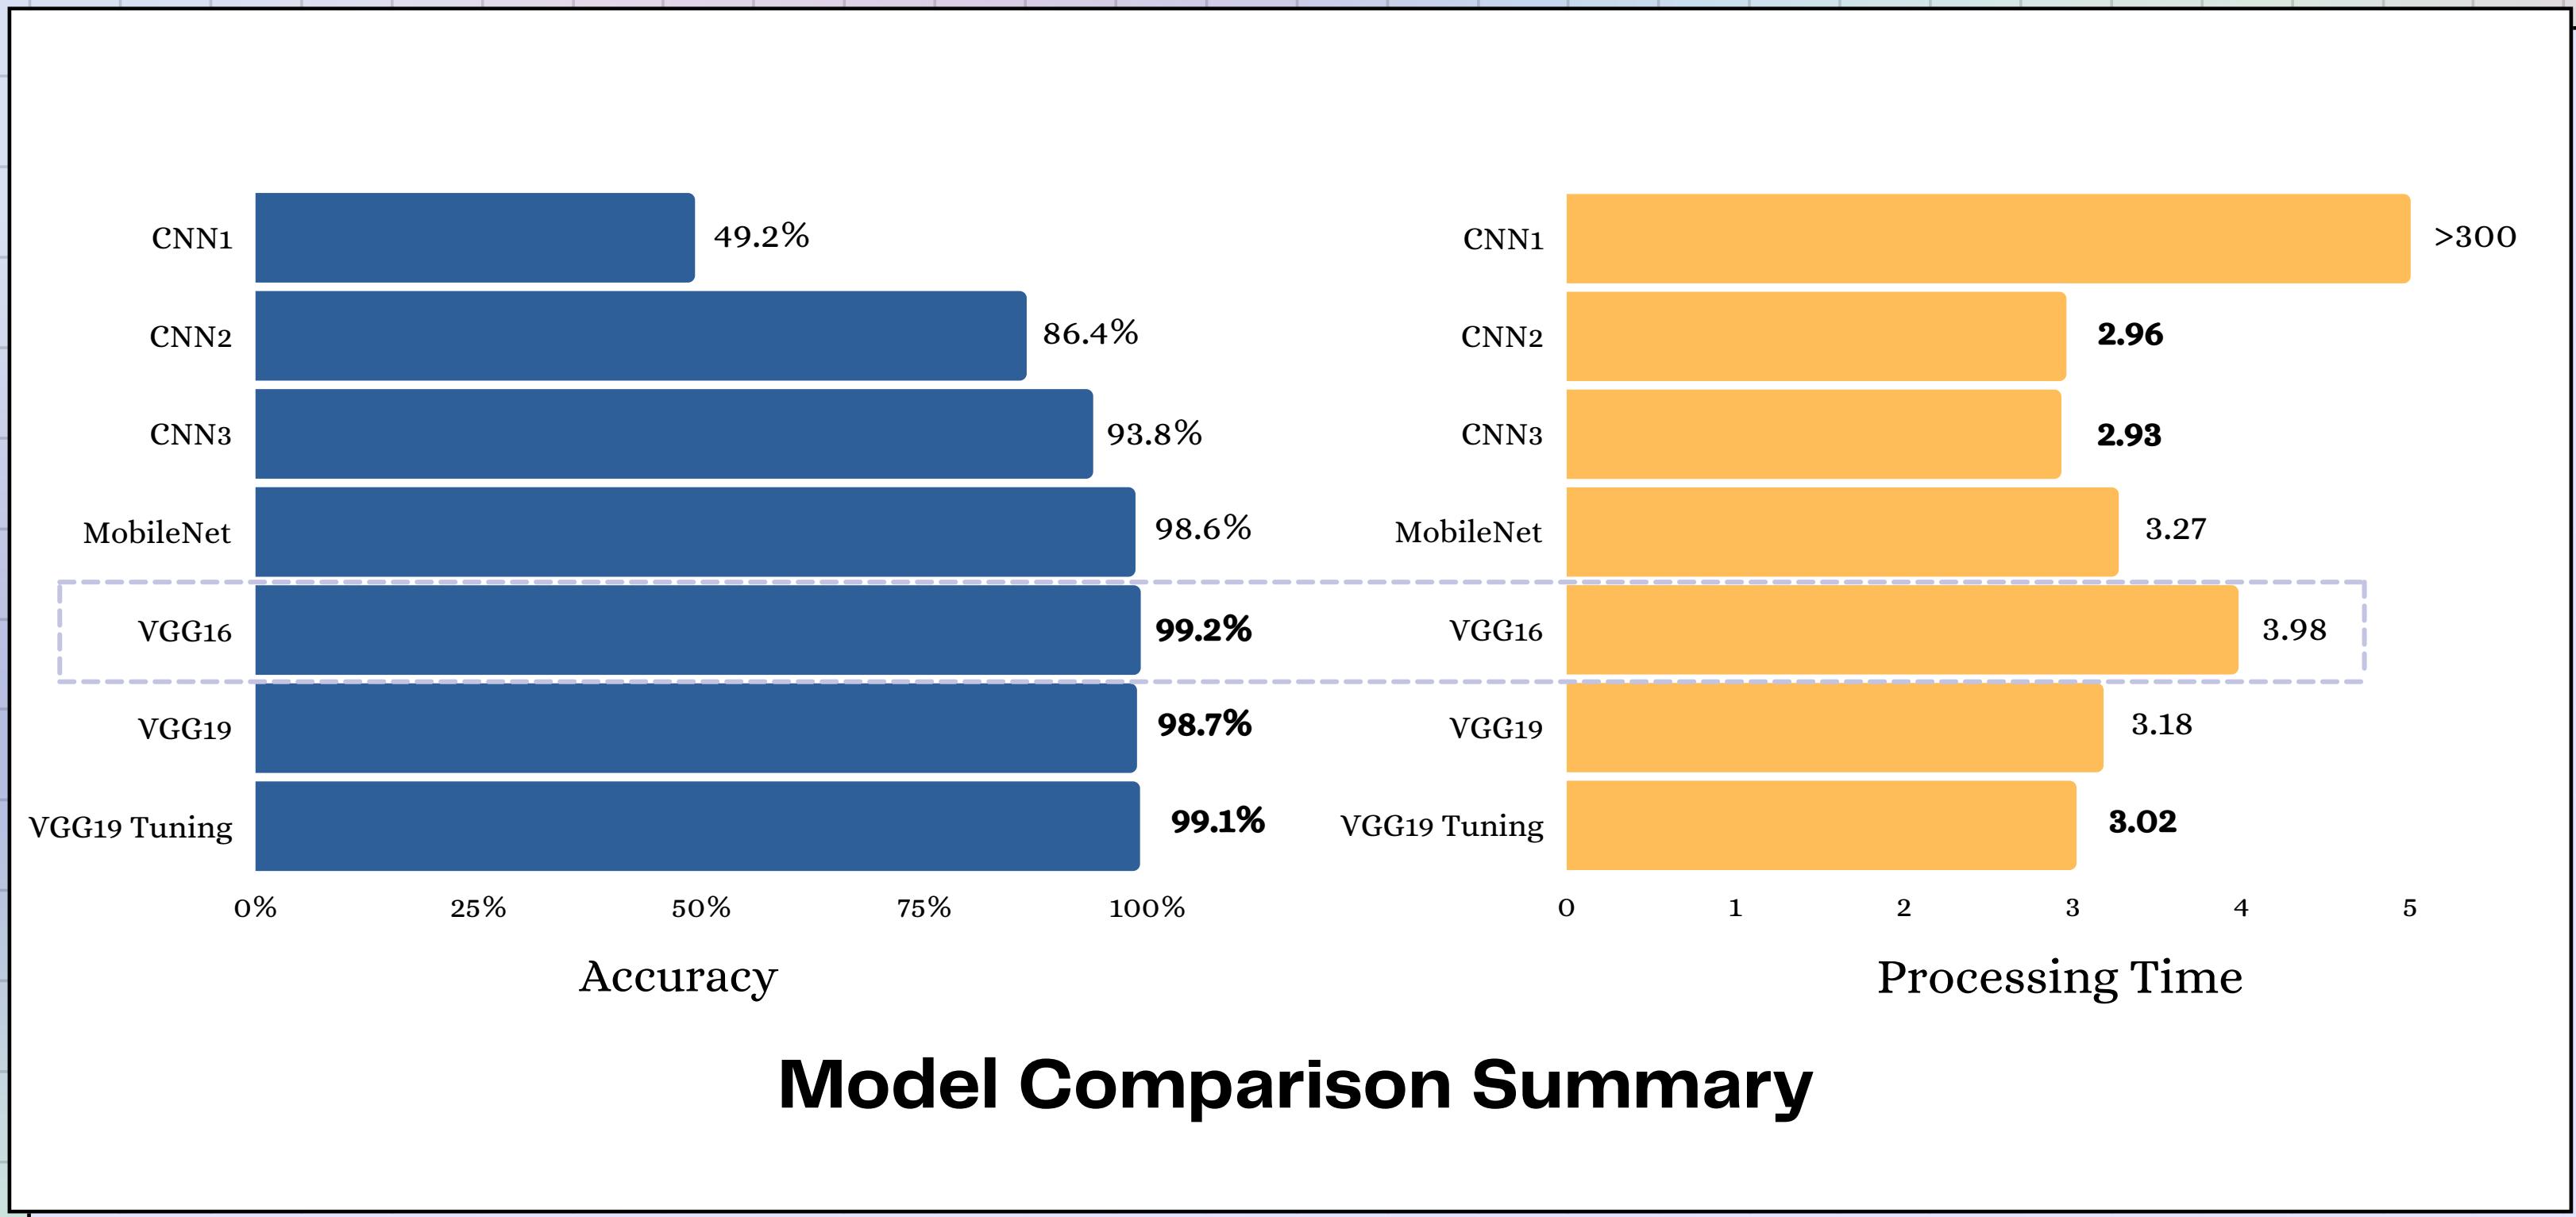
○ ○ ○ ○

Test Accuracy: 99.1%



Sample Test & Mask Detection

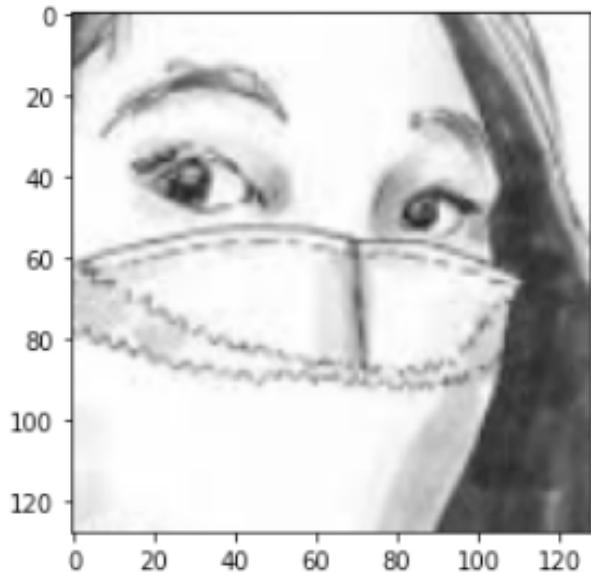
○ ○ ○ ○



Sample Test



WithMask Test

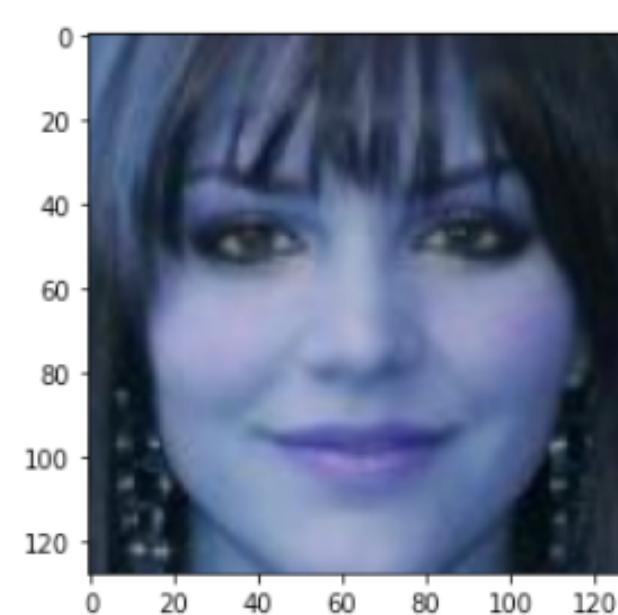


```
best_model.predict(sample_mask_img)  
array([[9.9984312e-01, 2.4890372e-07]])
```

```
[ ] pred = best_model.predict(sample_mask_img)  
predicted_class_indices = np.argmax(pred, axis=1)  
  
labels = (train_gen.class_indices)  
labels = dict((v,k) for k,v in labels.items())  
predictions = [labels[k] for k in predicted_class_indices]  
predictions  
  
['mask']
```



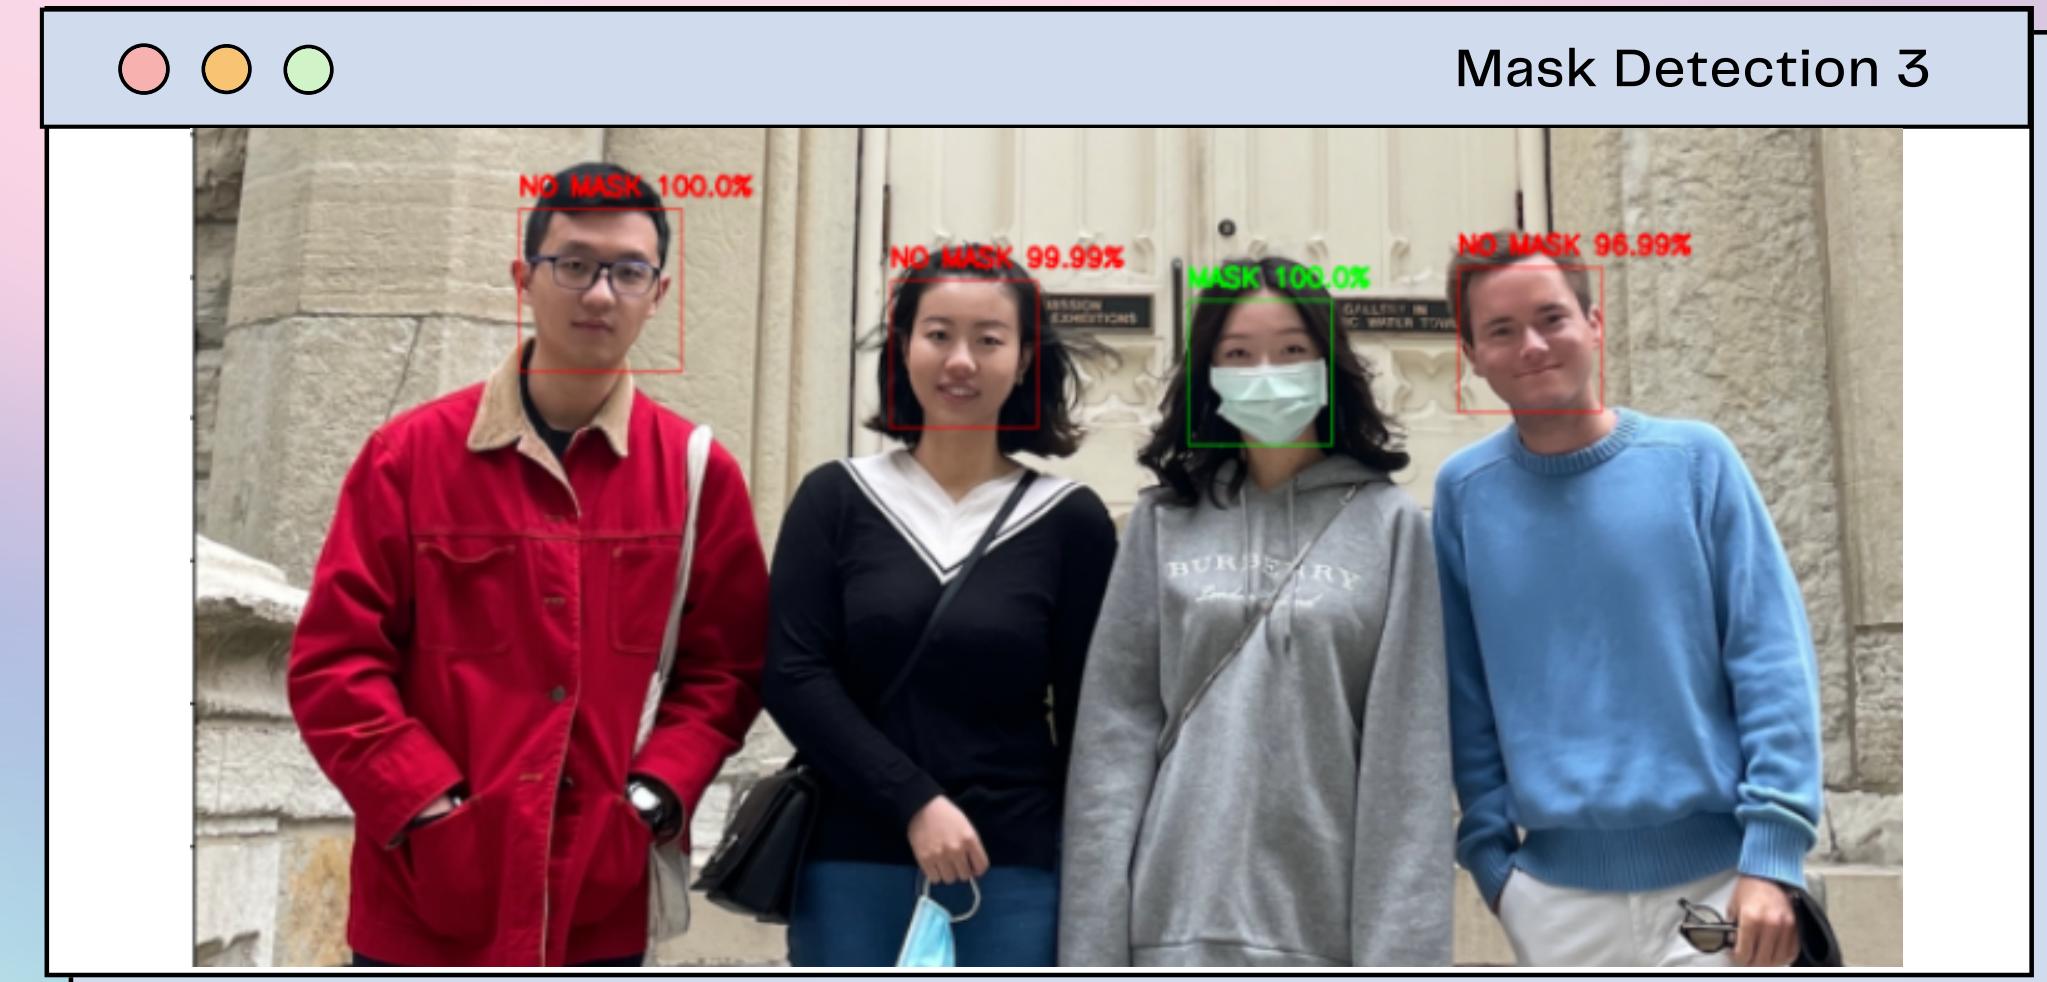
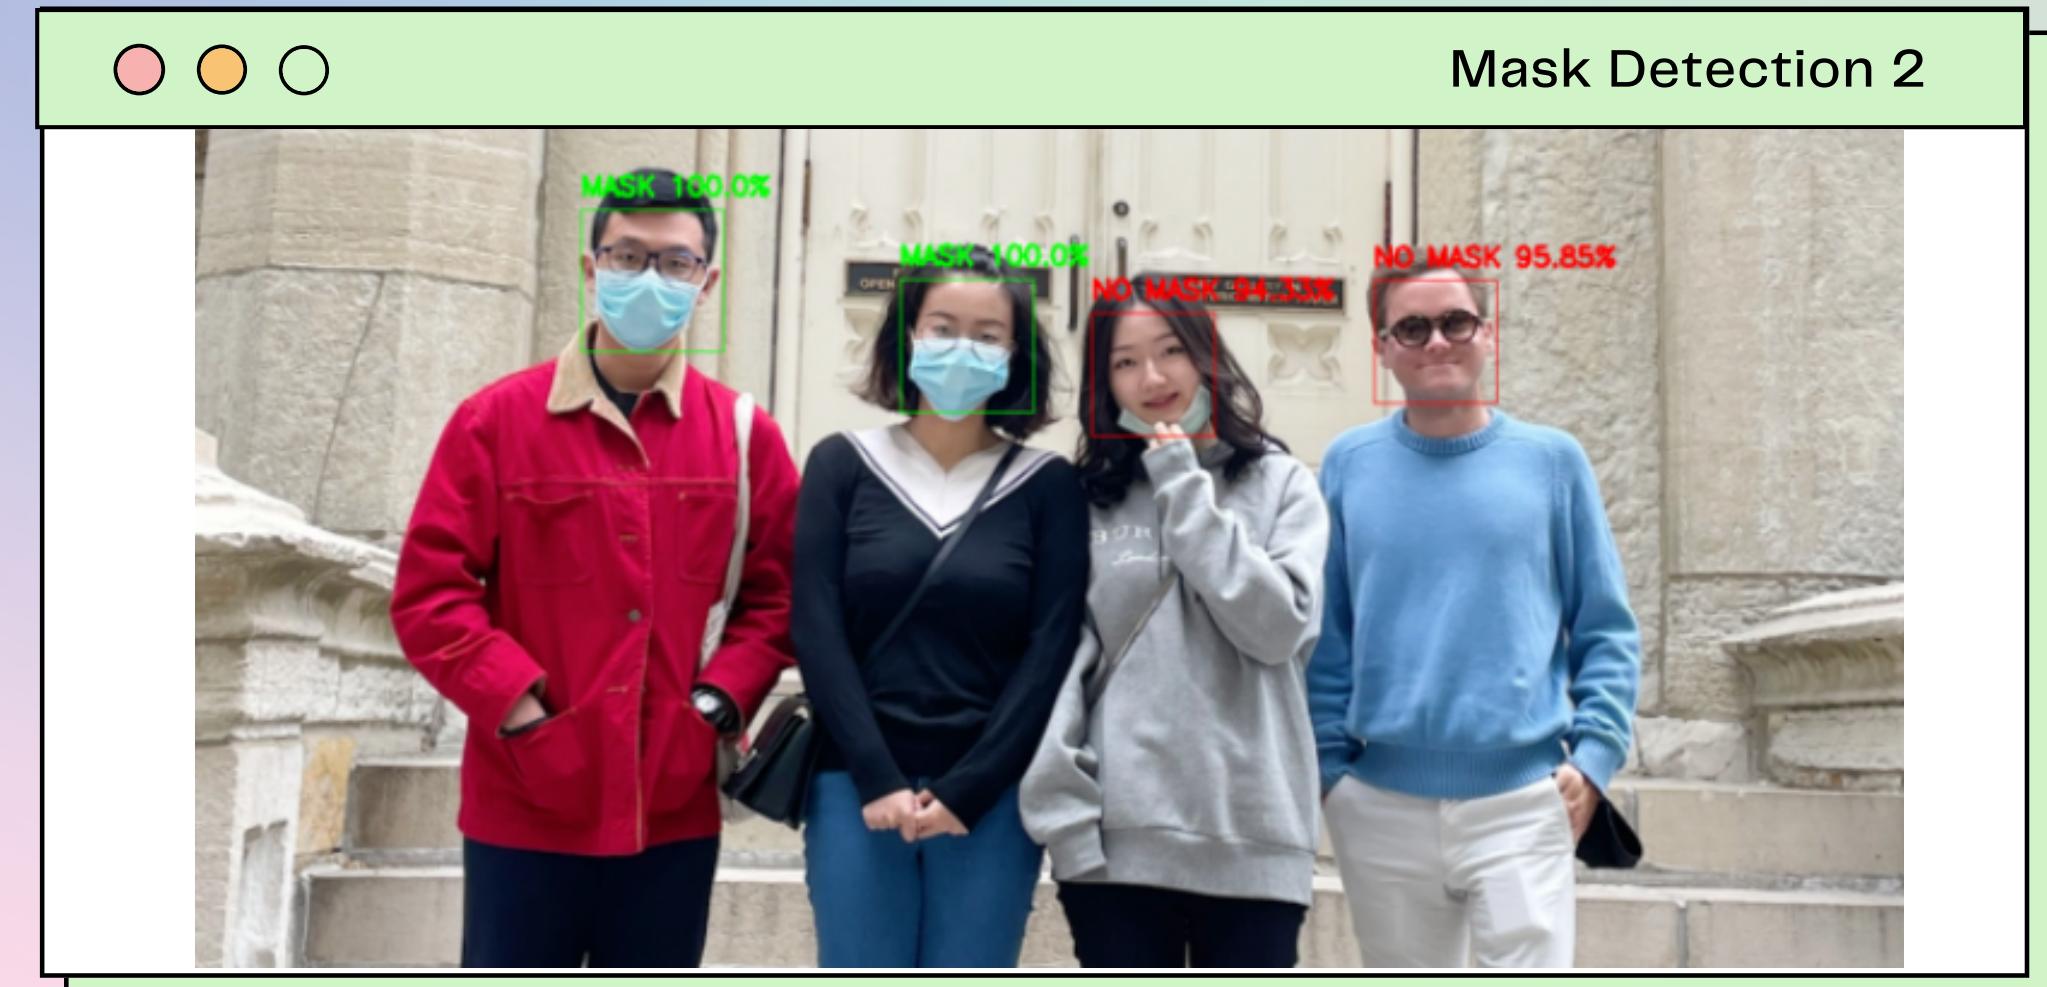
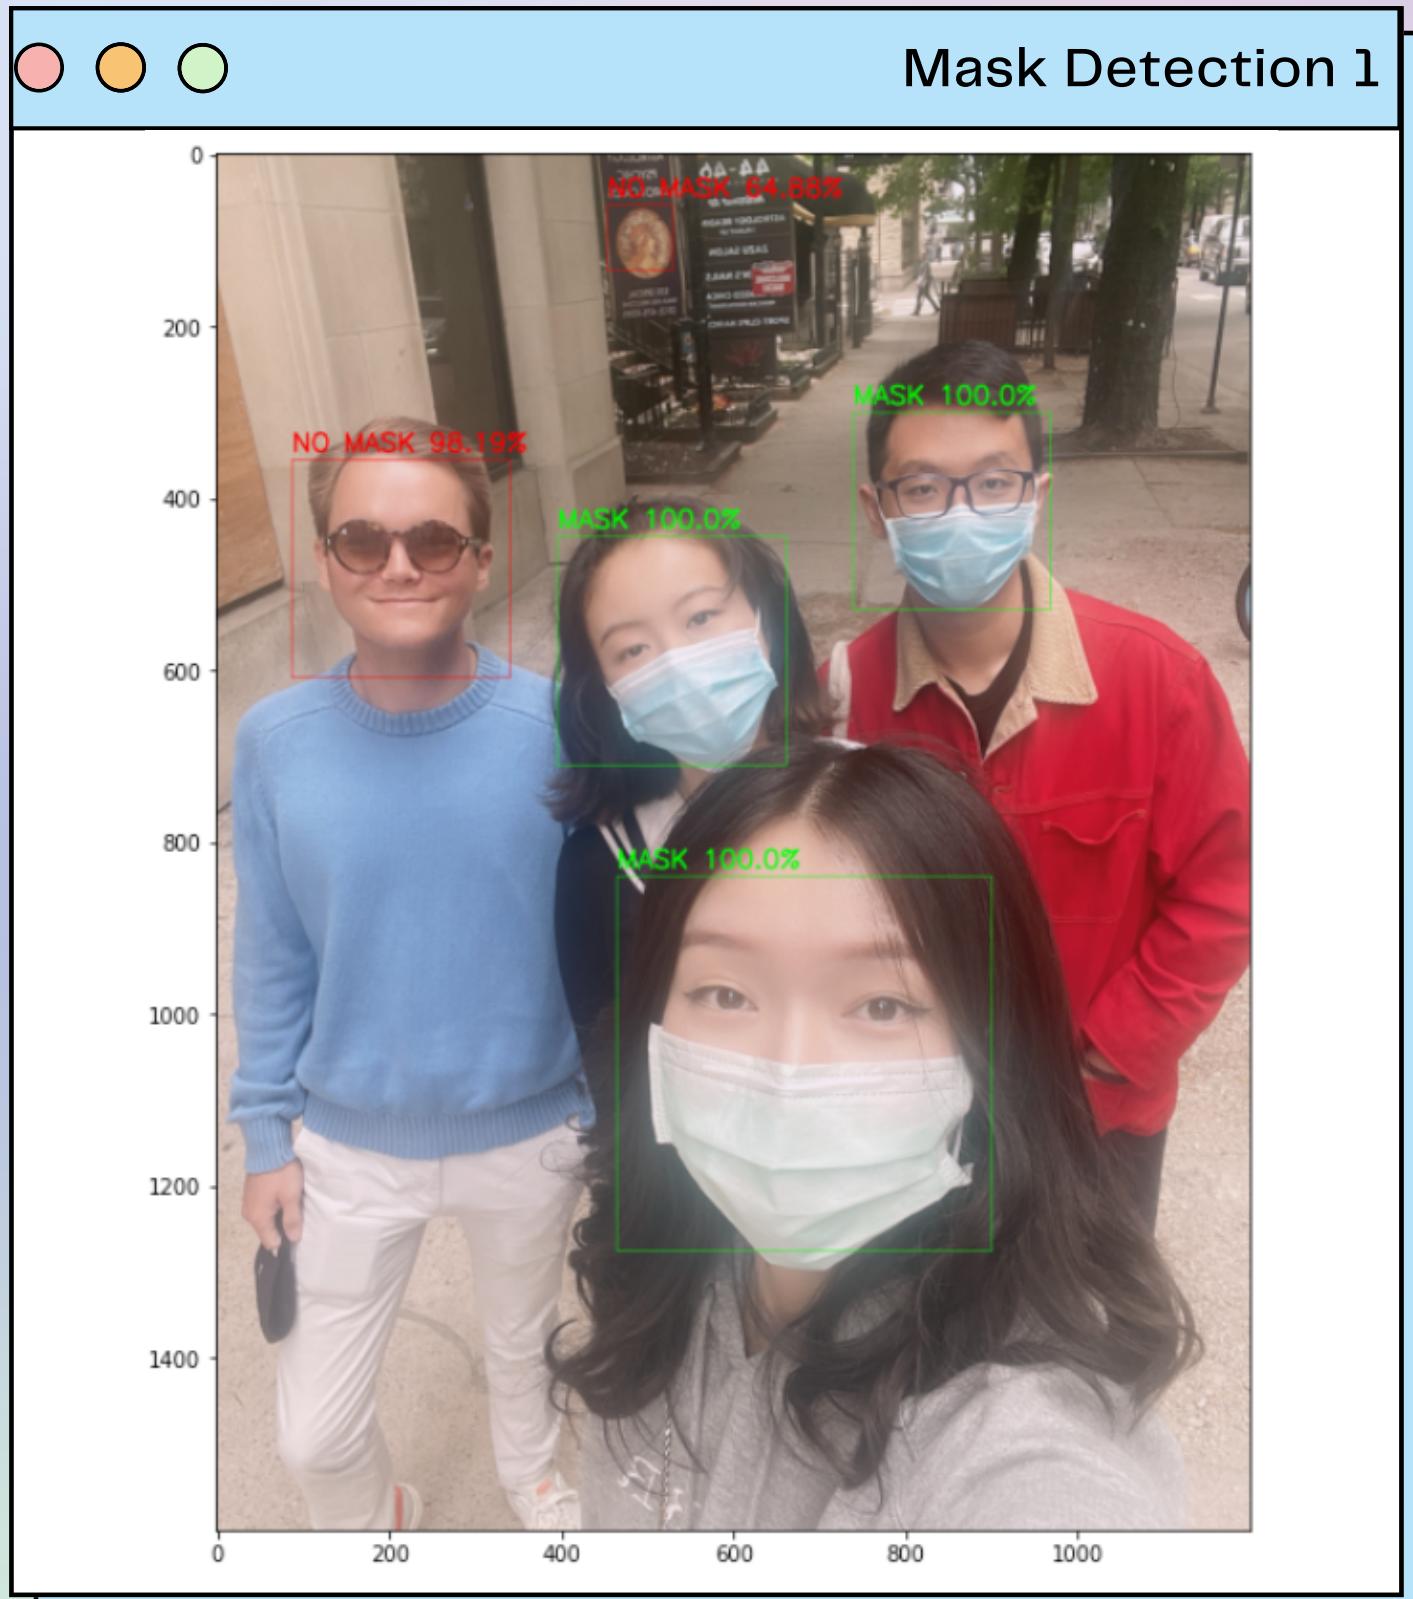
WithoutMast Test



```
best_model.predict(sample_no_mask_img)  
array([[1.2252026e-05, 9.9854195e-01]])
```

```
[ ] pred = best_model.predict(sample_no_mask_img)  
predicted_class_indices=np.argmax(pred, axis=1)  
  
labels = (train_gen.class_indices)  
labels = dict((v,k) for k,v in labels.items())  
predictions = [labels[k] for k in predicted_class_indices]  
predictions  
  
['no_mask']
```

Mask Detection Test



**That's a
wrap!**

o o o o

Thank you
for listening.

References

o o o o

Data

- <https://www.kaggle.com/andrewmvd/face-mask-detection>
- <https://www.kaggle.com/lalitharajesh/haarcascades>
- <https://www.kaggle.com/nageshsingh/mask-and-social-distancing-detection-using-vgg19>

Models

- <https://www.robots.ox.ac.uk/~vgg/>
- <https://ai.googleblog.com/2017/06/mobilennets-open-source-models-for.html>
- Andrew G. Howard, Menglong Zhu, Bo Chen, Dmitry Kalenichenko, Weijun Wang, Tobias Weyand, Marco Andreetto, & Hartwig Adam. (2017). MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications.